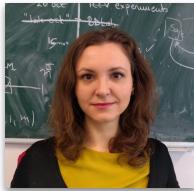

Deep Learning

— 1. Introduction to Deep Learning —

Course Team



**Florin
Brad**



**Elena
Burceanu**



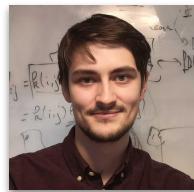
**Marius
Dragoi**



**Florin
Gogianu**



**Emanuela
Haller**



**Andrei
Manolache**



**Andrei
Nicolicioiu**

Invited Speaker



Iulia Dută

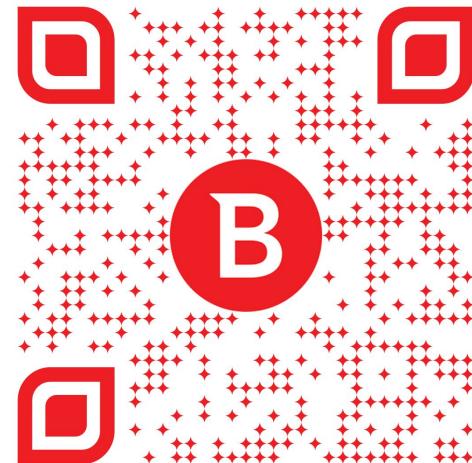
**University of
Cambridge**

{fbrad, eburceanu, mdragoi, fgogianu, ehaller, amanolache, anicolicioiu}@bitdefender.com

id366@cam.ac.uk

Our Team

Team page



bit-ml.github.io

AI / ML Community

Bucharest

- [Bucharest Deep Learning](#)
- [Bucharest AI](#)

Events:

- [Romanian AI Days](#)
- [EEML Summer School](#)

AI Romania

<https://www.airomania.eu/>



Eastern European Machine Learning Summer School

6-14 July 2022, Vilnius Lithuania
(hybrid event)

<https://www.eeml.eu/>

- **BONUS** points for acceptance

- Deep Learning lectures
- Tutorial sessions
- Paper reading sessions
- Poster sessions
- Top speakers from
**DeepMind, Google, McGill,
Mila, Cambridge, Harvard,
Oxford, Microsoft, Samsung**
- Application
 - Extended abstract
 - Deadline: 7 April 2022

AI / ML Community

Research Centers:

FMI: Marius Popescu, Bogdan Alexe, Radu Ionescu et al.

UPB: Traian Rebedea, Marius Leordeanu et al.

IMAR: Marius Leordeanu et al.

UT (Cluj): Lucian Bușoniu et al.

Bitdefender: bit-ml.github.com team

Siemens (Brașov): Lucian Itu et al.

ROVIS (Brașov): Sorin Grigorescu et al.

Industry:

Bitdefender: ML

Arnia: Outsourcing CV, NLP

Sparktech: Outsourcing CV, NLP

Xperi (FotoNation): CV

Siemens: Medical CV, ML

BOSCH (Cluj): Autonomous driving

Continental (Timișoara): Autonomous driving

Administrative

Program

- **Lectures**
 - Each Wednesday: 18:00 - 20:00
 - 3 speed tests
 - 5 minutes
 - 1-2 short questions
- **Labs**
 - Each Wednesday: 16:00 - 18:00
- **Course Materials**

https://drive.google.com/drive/folders/1I_ujnkkVxOSMxOa8IAuRp0NmaXzVraXF?usp=sharing

Assignments and Projects

- **2 Assignments (PyTorch)**

- Related to the lab content
- Submission deadline: 2 or 3 weeks after publishing
 - penalties for delays
- Both of them will be presented in Week 9

- **1 Project (PyTorch)**

- Teams of 2 members
- Subjects will be proposed in Week 4
- Each student should choose a subject until Week 6
- Will be presented in Week 10 - Poster Session

Course Structure

Week	Lecture	Speed Test	Laboratory	Event	Deadline
1	Introduction to Deep Learning		PyTorch and Tensors		
2	Neural Networks		Neural Networks		
3	Optimization		Training and Optimization	Assignment 1 out	
4	Debugging Neural Networks	Test 1	Debugging	Project proposals	
5	Convolutional Neural Networks		ConvNets (1)	Assignment 2 out	Assignment 1
6	Computer Vision Applications		ConvNets (2)		Project selection
7	Recurrent Neural Networks	Test 2	RNN		
8	Natural Language Processing Applications		NLP Applications		Assignment 2 EEML - April 7
9	Unsupervised Learning	Test 3	Assignments Evaluation		
10	Project Evaluation (Poster Session)		Project Evaluation		Project

Grading

	Requirements	
Assignment 1	2 points	
Assignment 2	2 points	
3 x Speed Tests	1 point	50%
[BONUS] Good answers / questions during lectures and labs	max 2 points	
Project	5 points	
[BONUS] Good interventions / questions during the poster session	max 1 point	50%

Deep Learning Course Structure



Lecture Overview





I. Deep Learning Motivation

Artificial Intelligence

ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



MACHINE LEARNING

Ability to learn without explicitly being programmed

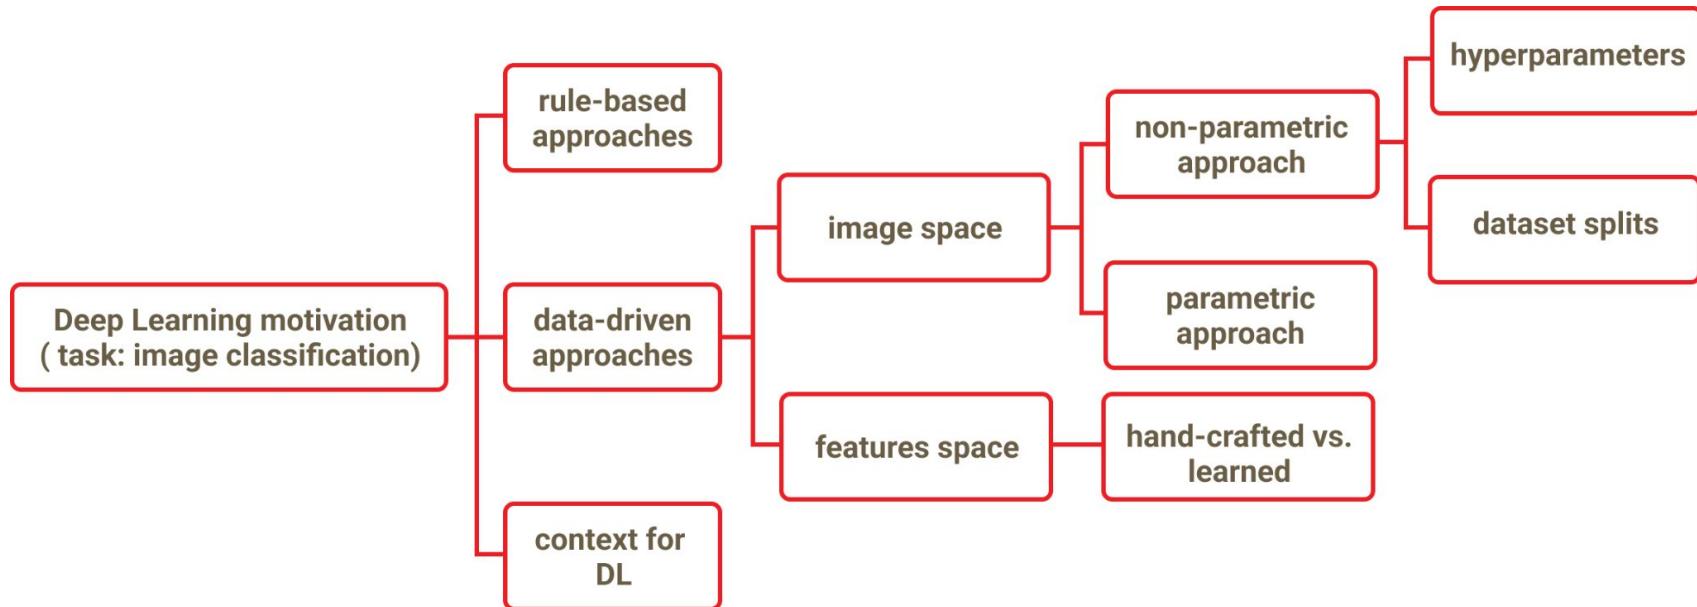


DEEP LEARNING

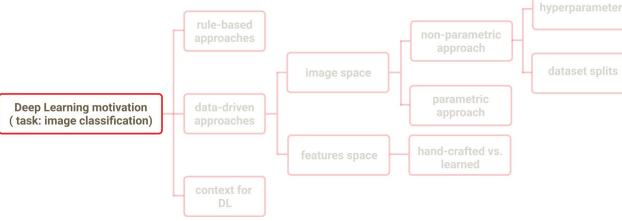
Extract patterns from data using neural networks



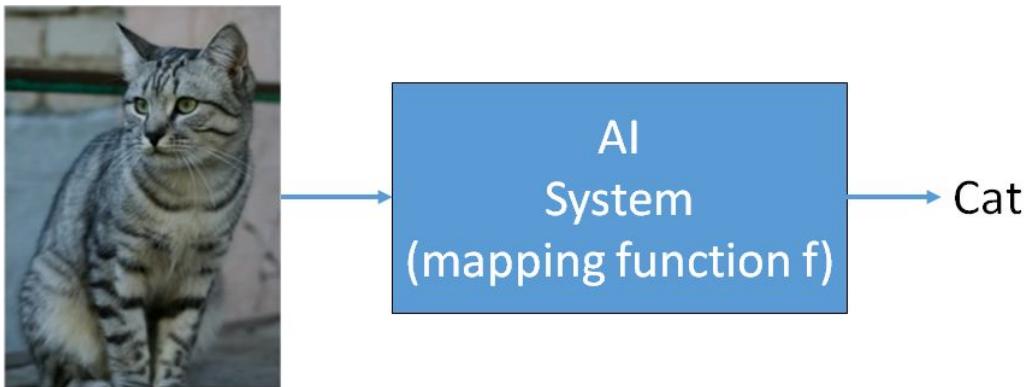
Deep Learning Motivation



Task: Image Classification

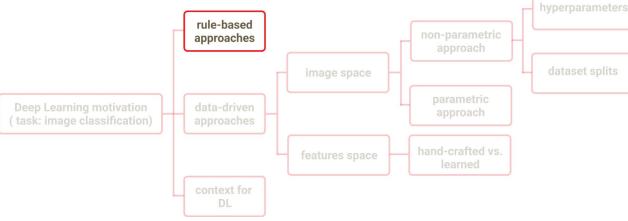
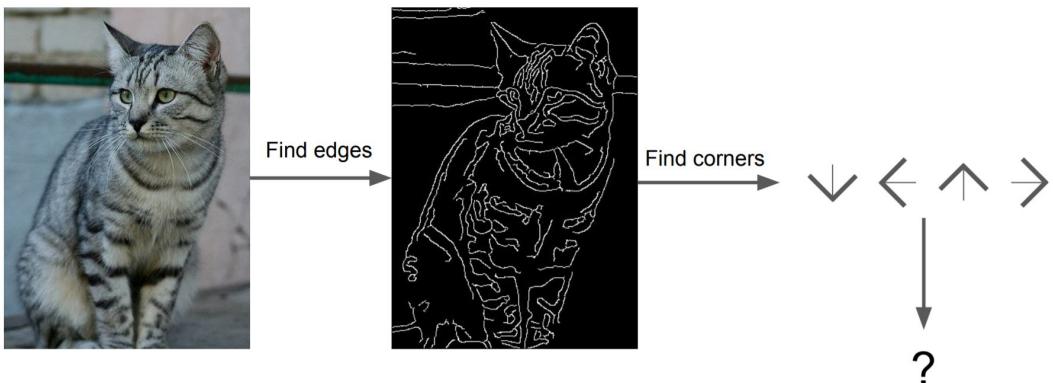


- Classify an image according to its visual content
- Consider a discrete set of labels
 - e.g. {dog, airplane, car, cat}

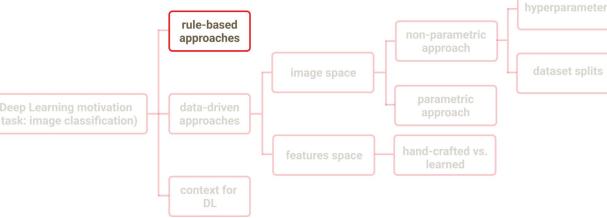


Rule-Based Approach

- Hand crafted features and rules



Challenges



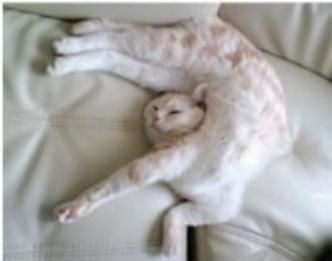
Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter

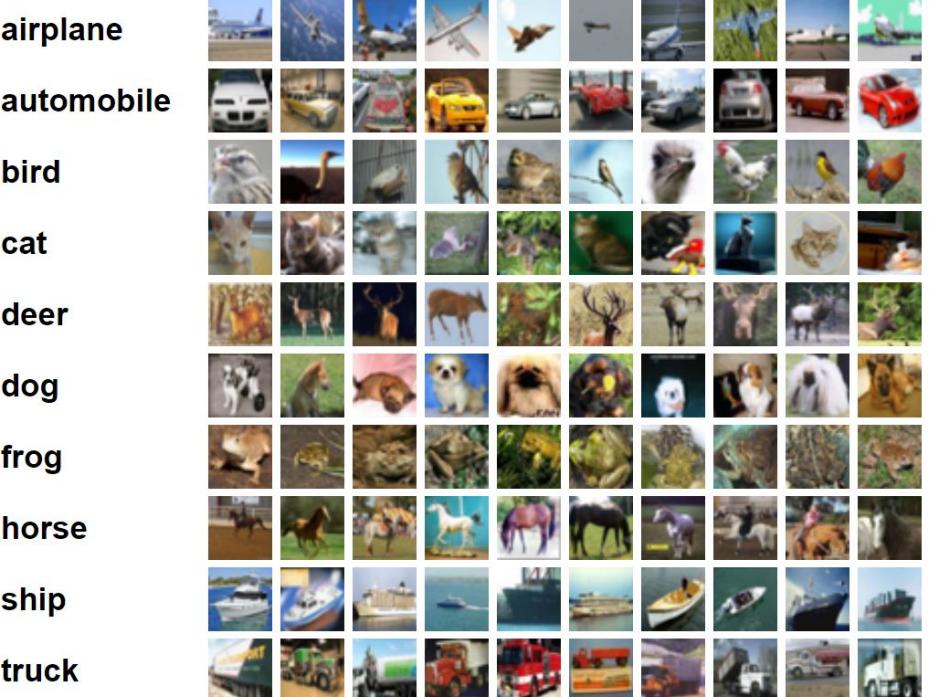
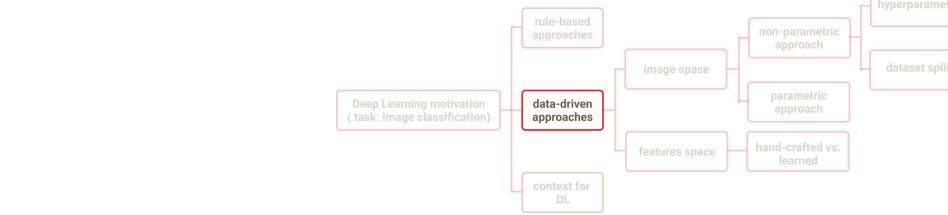


Intra-class variation



Data-Driven Approach

- Dataset
 - with images and associated labels
- CIFAR-10
 - 10 classes
 - 60 000 examples
 - $32 \times 32 \times 3$



CIFAR-10

Image Space

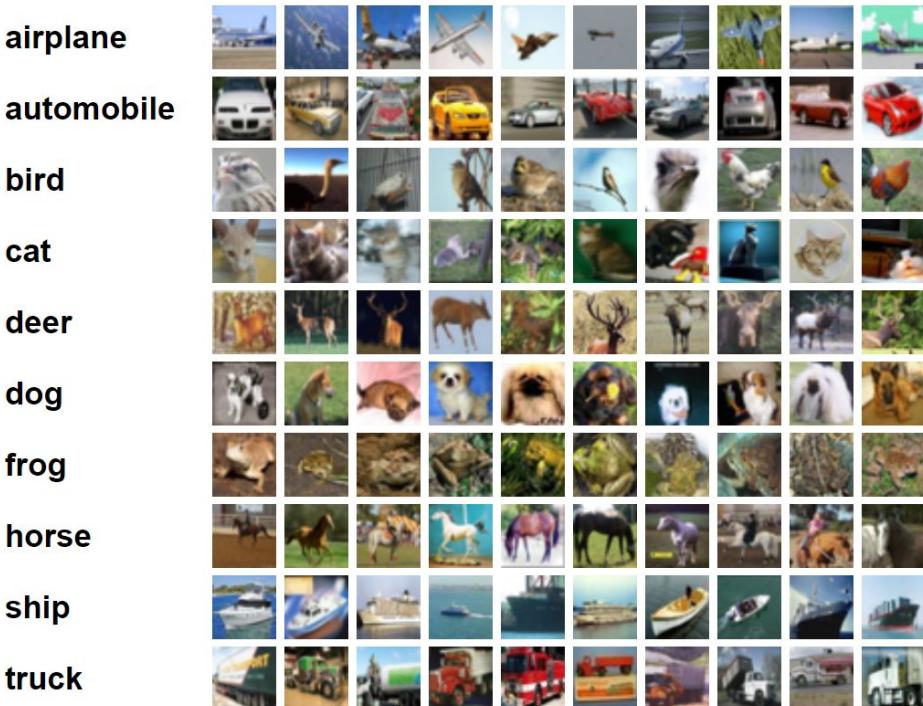
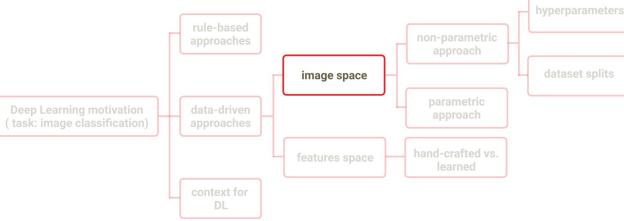
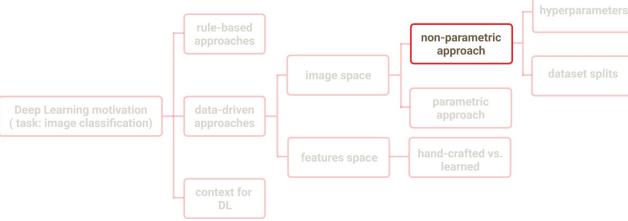
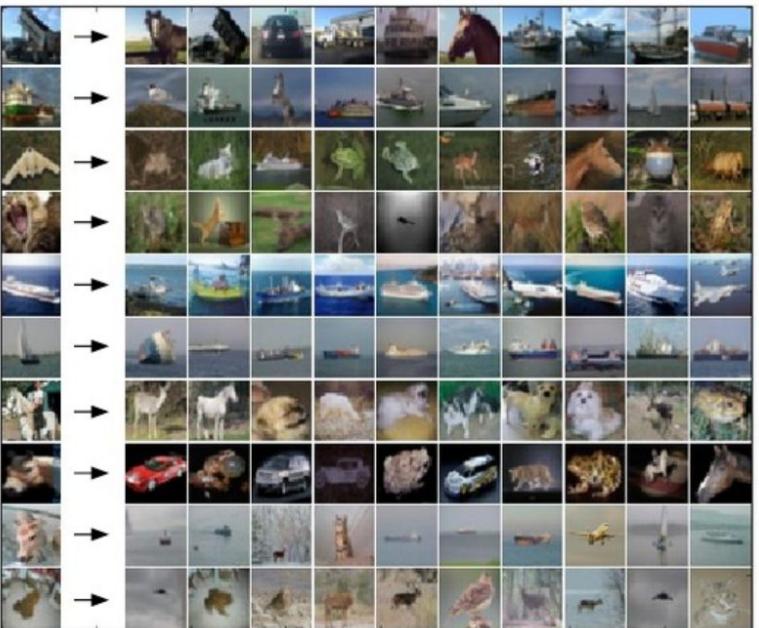


Image Space

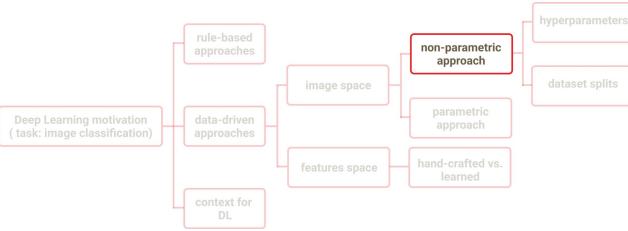
Non-Parametric Approach

K-Nearest Neighbor Classifier

- Memorize all examples and corresponding labels
- Predict label of top K most similar examples
- Define distance between two images



Distance metric

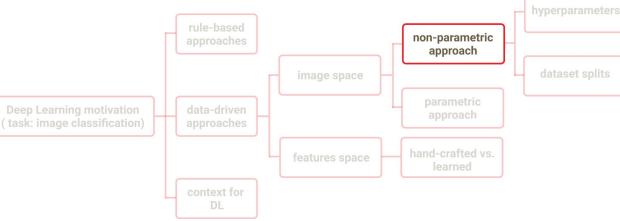
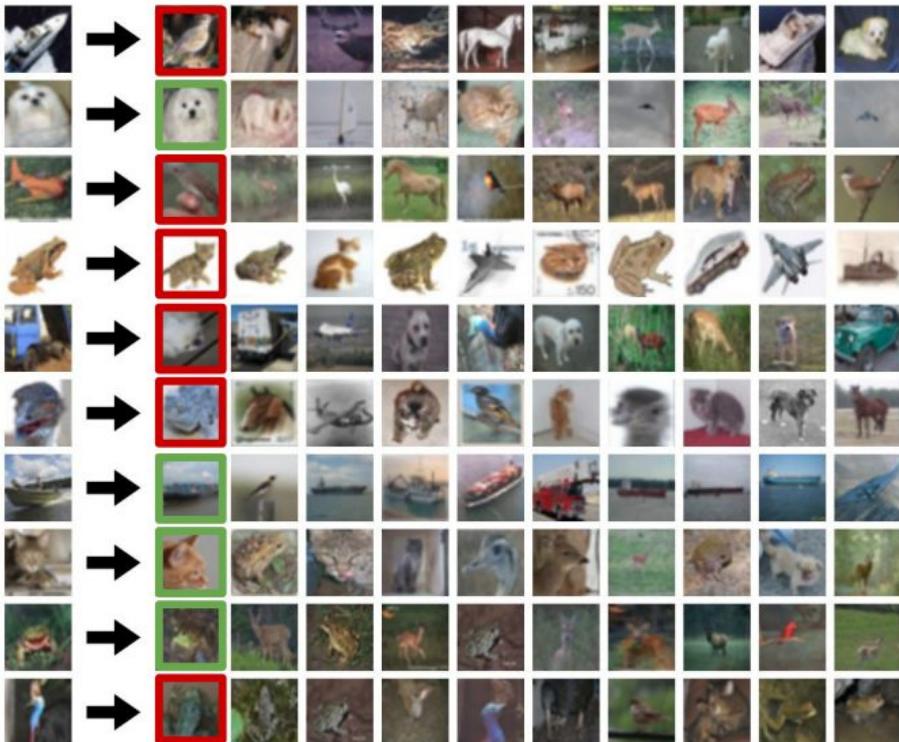


L1 distance: $d(I_1, I_2) = \sum_x |I_{1;x} - I_{2;x}|$, where $I_{1;x}$ is the x 'th element of image I_1

test image				training image				pixel-wise absolute value differences				→ 456
56	32	10	18	-	10	20	24	17	46	12	14	
90	23	128	133		8	10	89	100	82	13	39	33
24	26	178	200		12	16	178	170	12	10	0	30
2	0	255	220		4	32	233	112	2	32	22	108

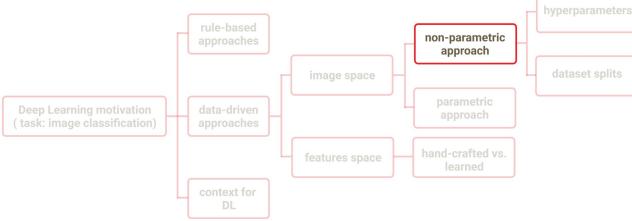
K-Nearest Neighbor Classifier

- K=1



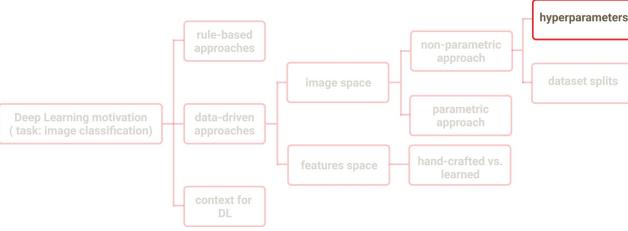
K-Nearest Neighbor Classifier

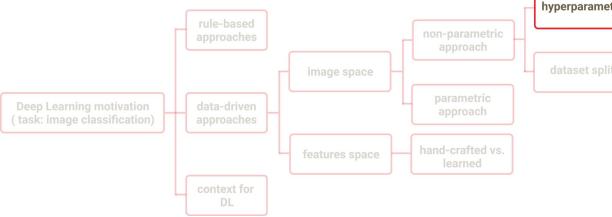
- Decision regions
 - $K = 1$



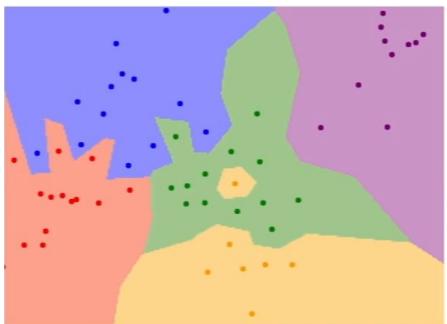
K-Nearest Neighbor Classifier

- Hyperparameters
 - K
 - Distance metric





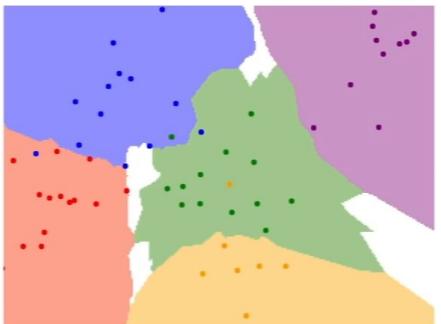
Hyperparameters: K



$K = 1$

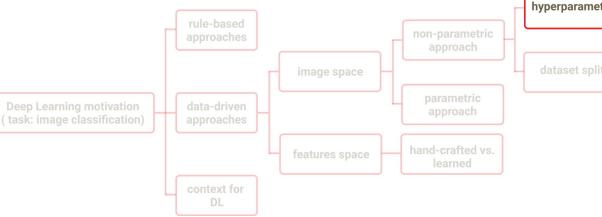


$K = 3$



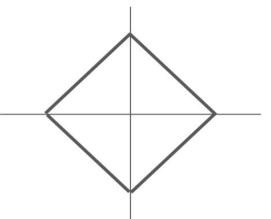
$K = 5$

Hyperparameter: distance metric



L1 (Manhattan) distance

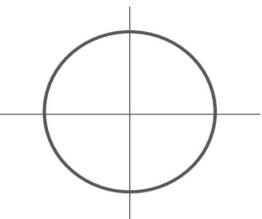
$$d_1(I_1, I_2) = \sum_x |I_{1;x} - I_{2;x}|$$



$K = 1$

L2 (Euclidean) distance

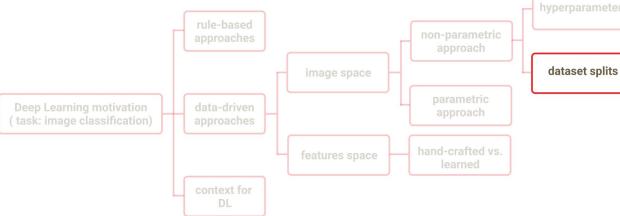
$$d_2(I_1, I_2) = \sqrt{\sum_x (I_{1;x} - I_{2;x})^2}$$



$K = 1$

Dataset splits

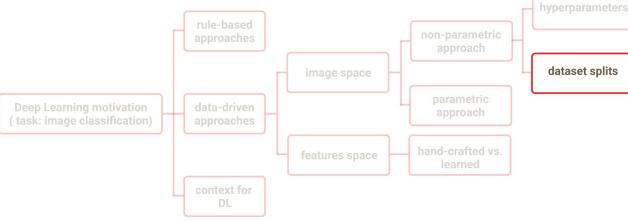
- Only train data



Train data

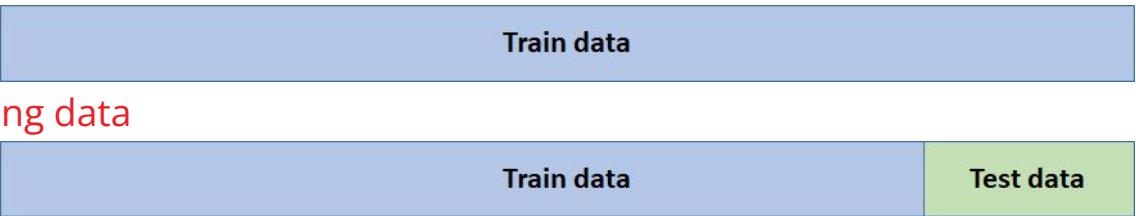
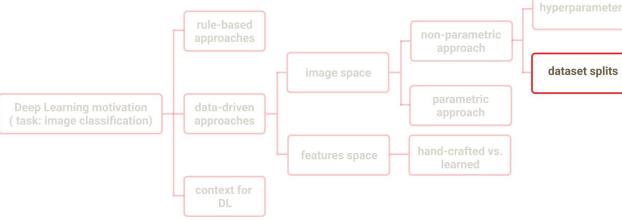
Dataset splits

- Only train data
 - K=1 - works best for training data

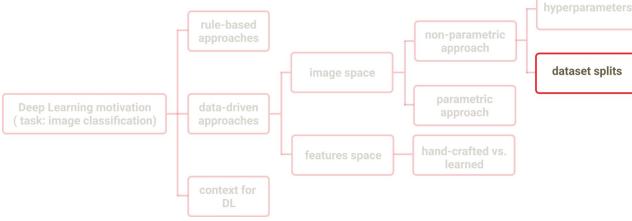


Dataset splits

- Only train data
 - K=1 - works best for training data
- Train data & Test data



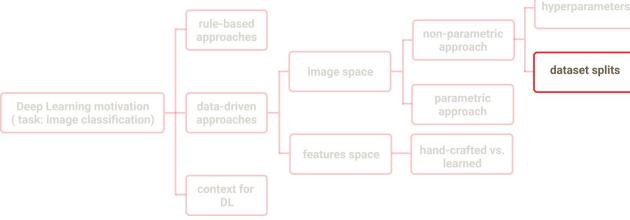
Dataset splits



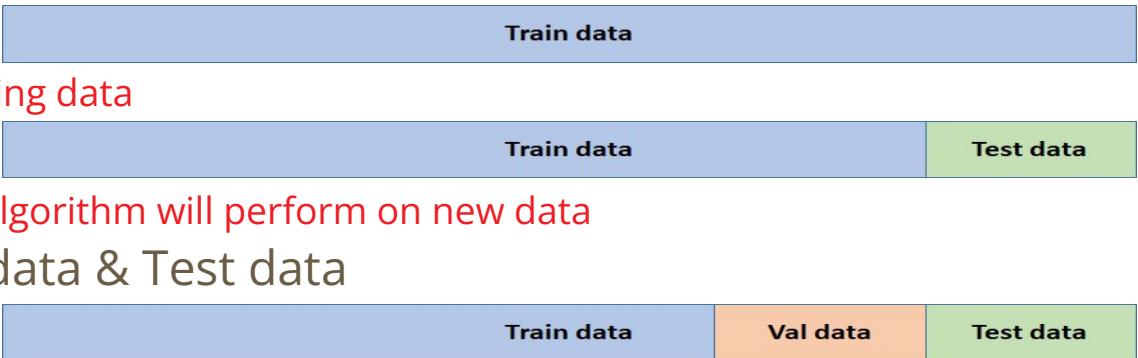
- Only train data
 - K=1 - works best for training data
- Train data & Test data
 - No intuition on how the algorithm will perform on new data



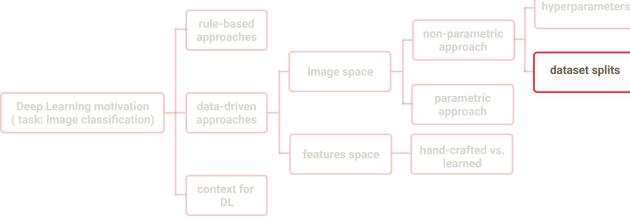
Dataset splits



- Only train data
 - K=1 - works best for training data
- Train data & Test data
 - No intuition on how the algorithm will perform on new data
- Train data & Validation data & Test data



Dataset splits



- Only train data



- K=1 - works best for training data

- Train data & Test data



- No intuition on how the algorithm will perform on new data

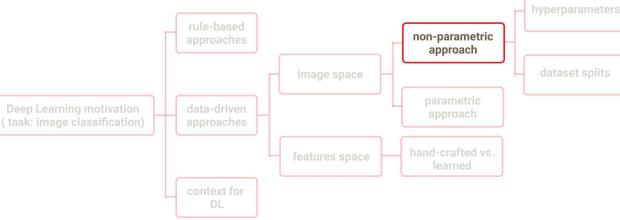
- Train data & Validation data & Test data



- Cross Validation

Fold 1	Fold 2	Fold 3	Fold 4	Test data
Fold 1	Fold 2	Fold 3	Fold 4	Test data
Fold 1	Fold 2	Fold 3	Fold 4	Test data
Fold 1	Fold 2	Fold 3	Fold 4	Test data

Image Space K-Nearest Neighbor



- Issues

- Pixel-wise distances are not informative
- Train: $O(1)$
- Test: $O(n)$

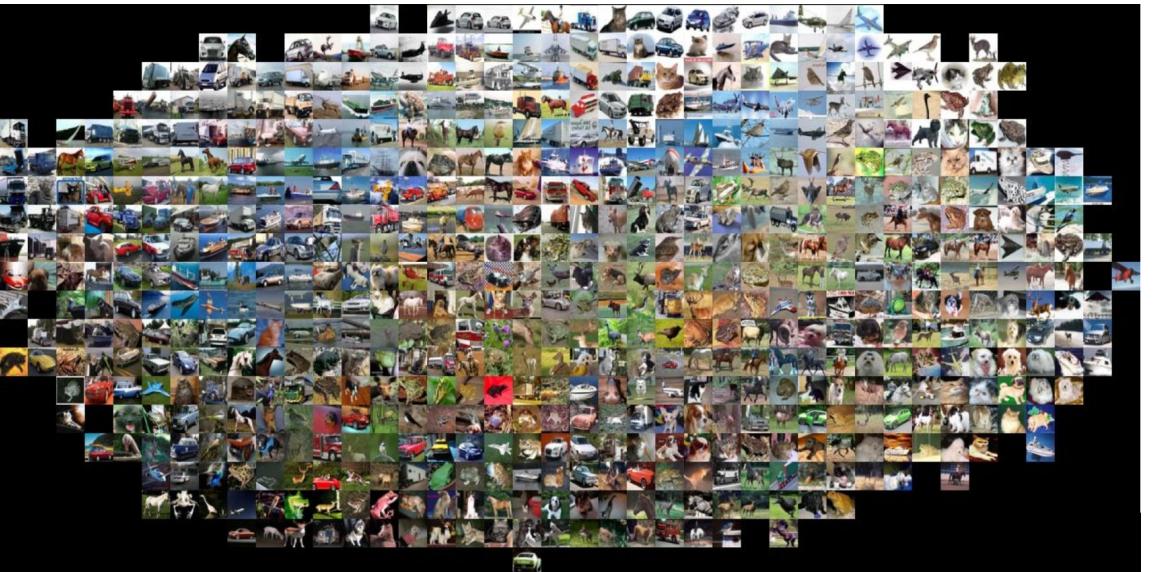
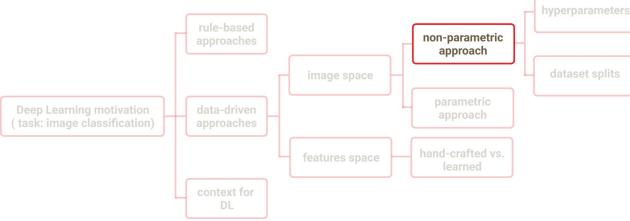
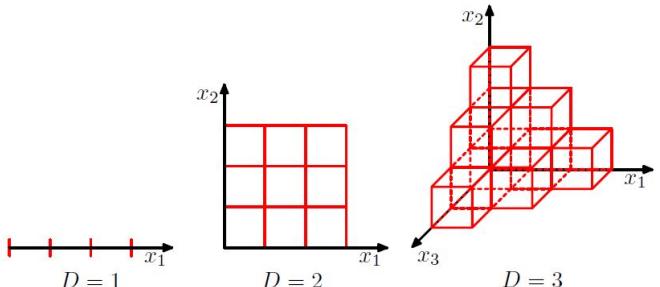


Image Space K-Nearest Neighbor

- Issues
 - Curse of dimensionality



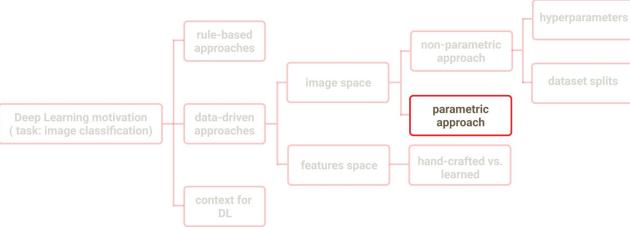
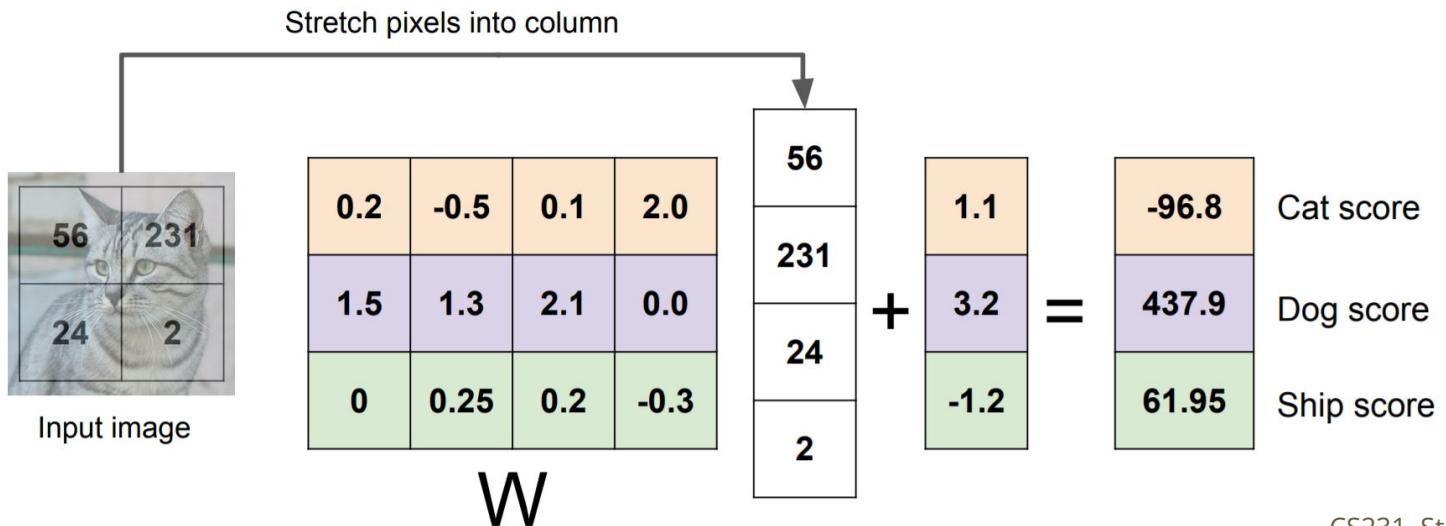


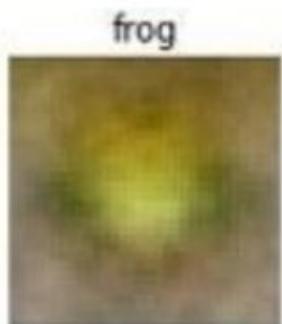
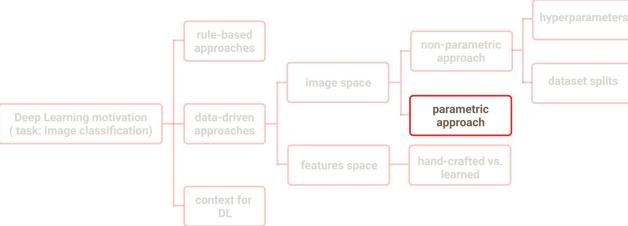
Image Space

Parametric approach - Linear Classifier

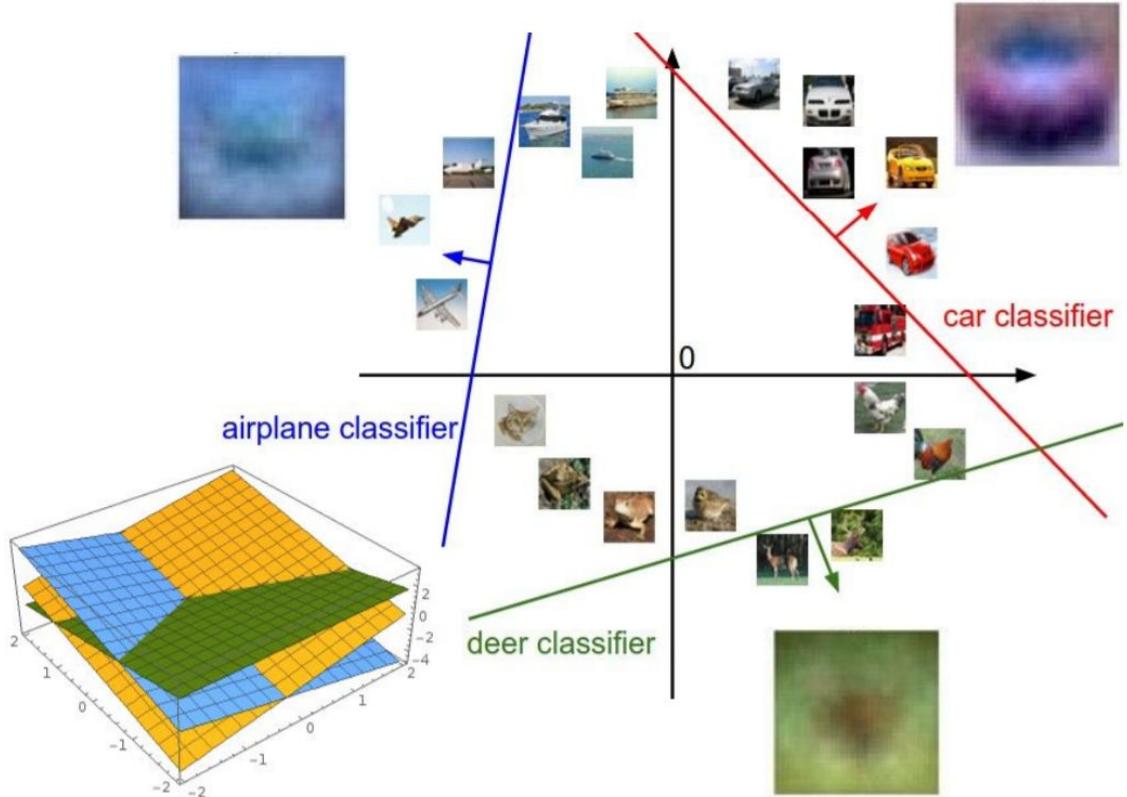
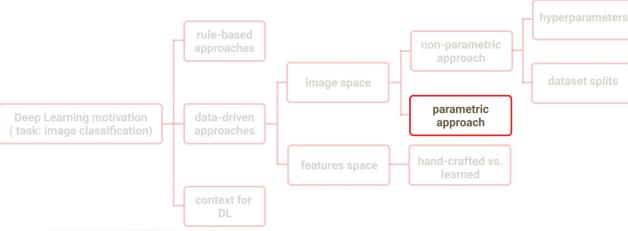
- $f(x, W) = Wx + b$, W - parameters to be learned



Linear Classifier - Interpretation



Linear Classifier - Interpretation



Features Space - Motivation

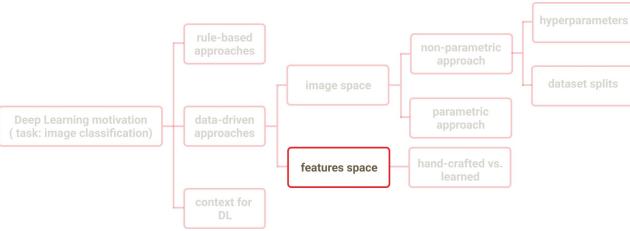
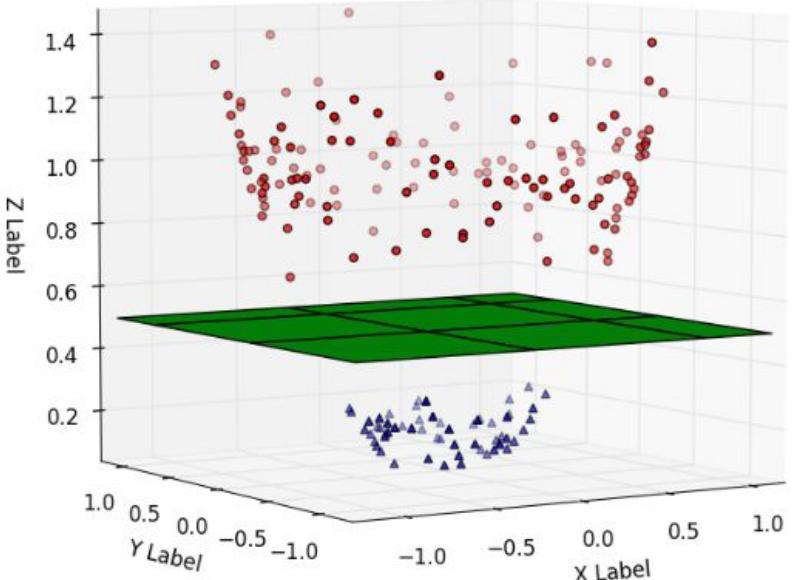
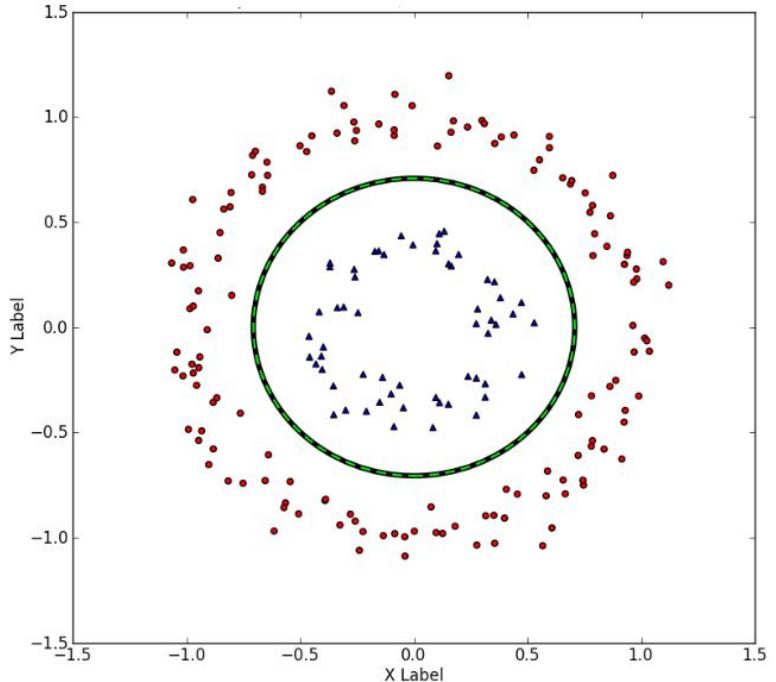
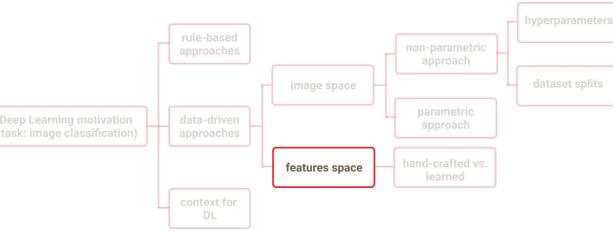
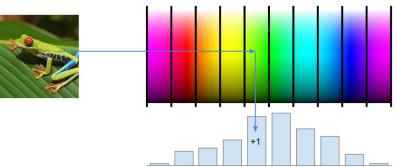


Image Space => Features Space

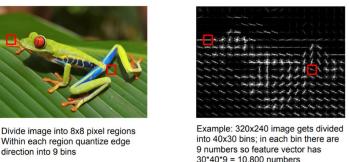


- Hand-crafted

Example: Color Histogram



Example: Histogram of Oriented Gradients (HoG)



Example: Bag of Words

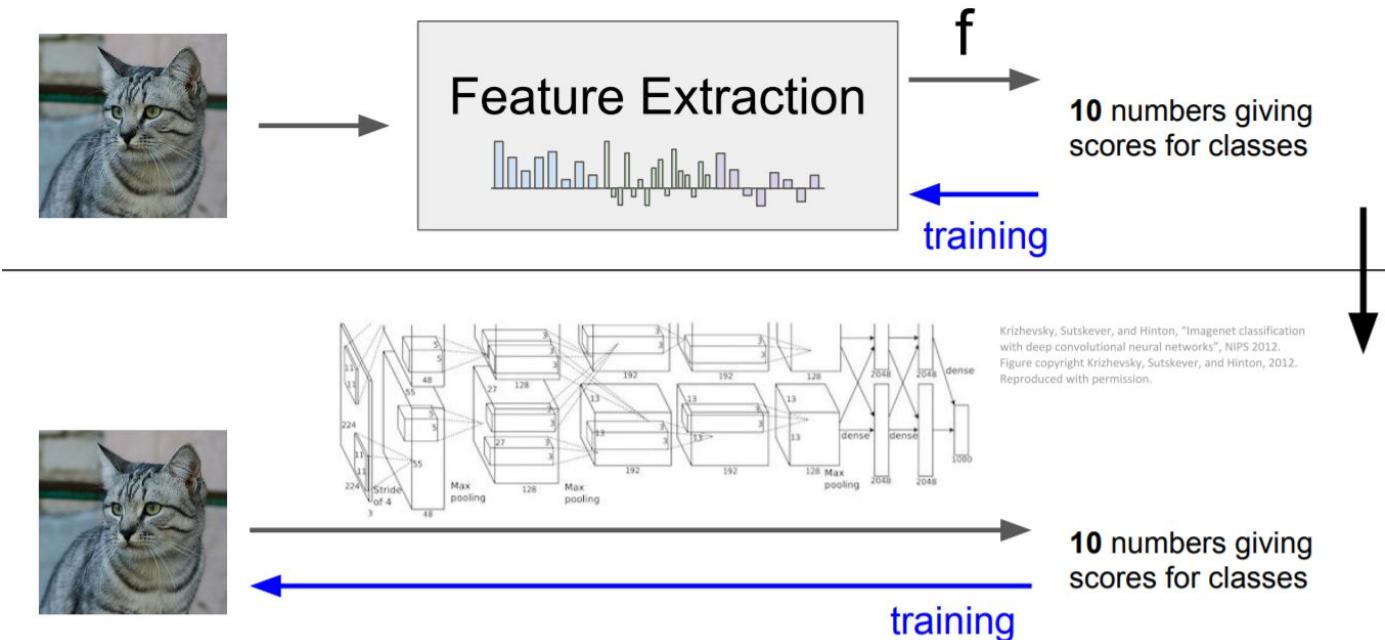
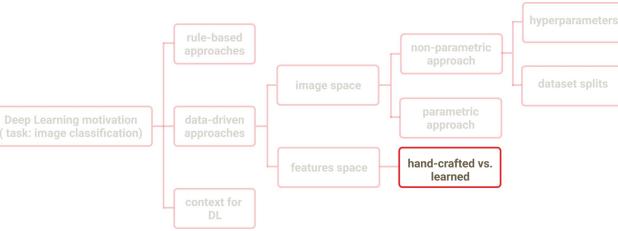
Step 1: Build codebook



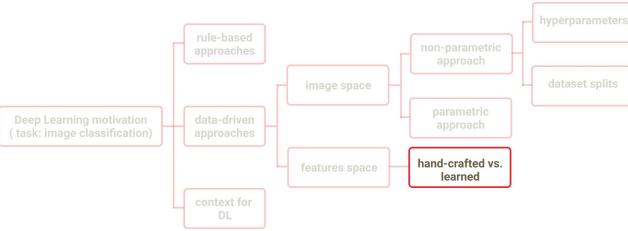
Step 2: Encode images



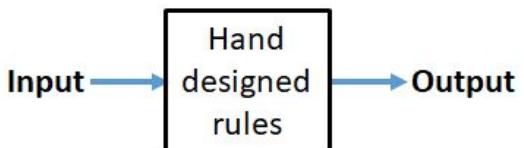
Hand-Crafted vs. Learned Features



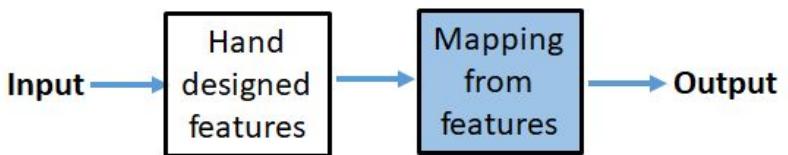
Hand-Crafted vs. Learned Features



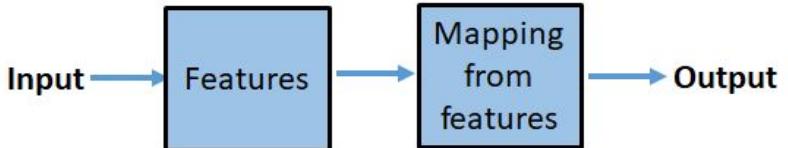
Rule-based systems



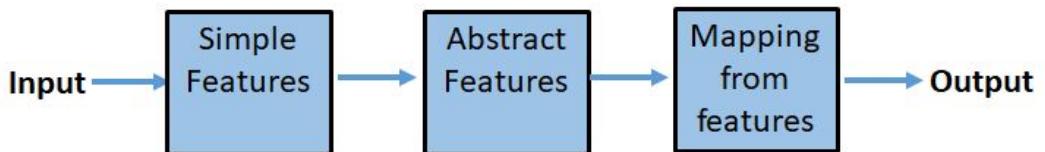
Classic ML



Representation learning

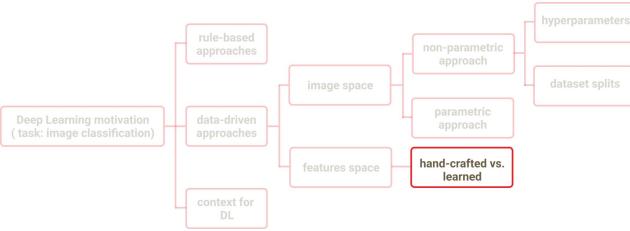


DL

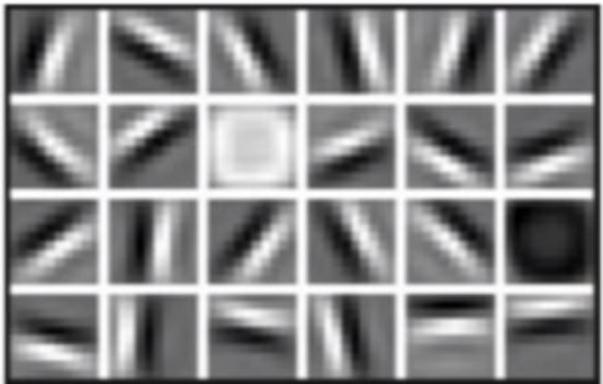


Why Deep Learning?

- Learn underlying features directly from data



Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

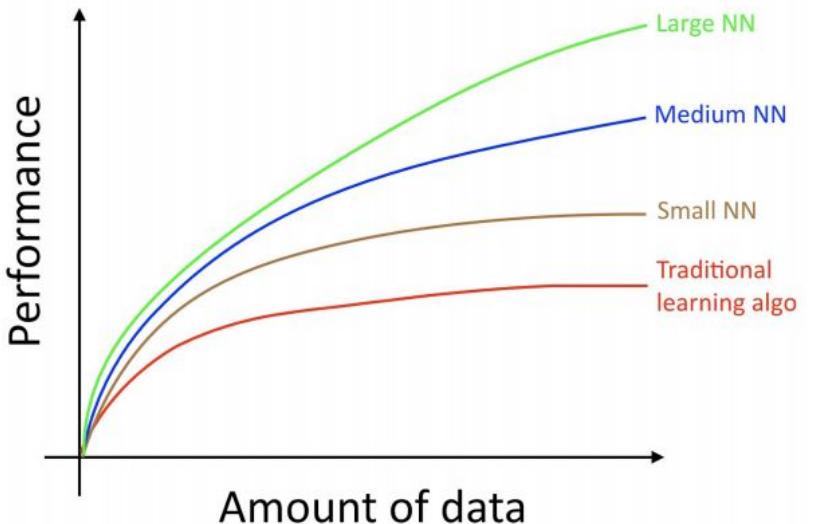
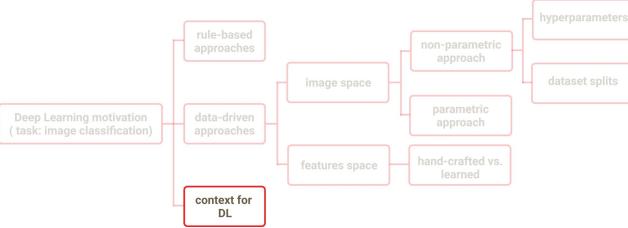
High Level Features



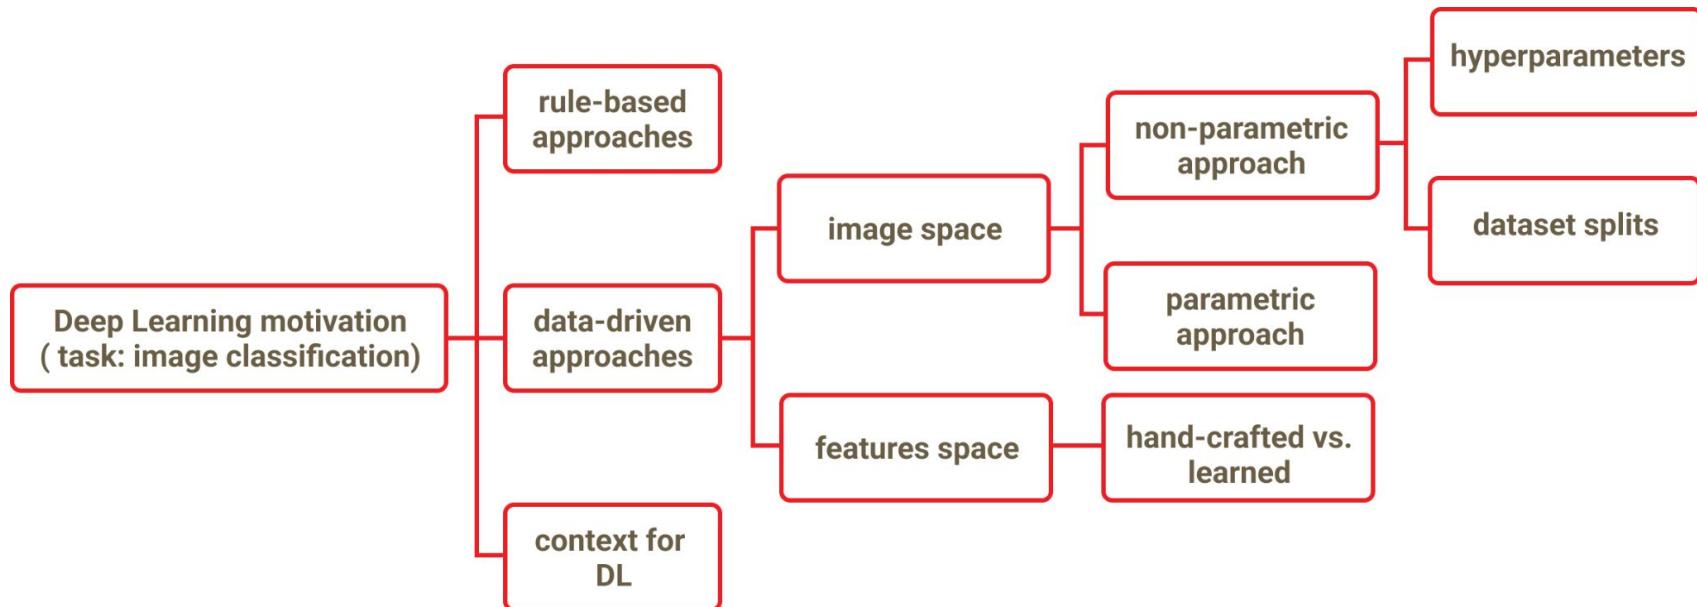
Facial Structure

Context for Deep Learning

- Big data
 - Larger datasets
- Hardware
 - GPUs
 - Massively parallelizable algorithms
- Deep Learning Frameworks
 - PyTorch
 - TensorFlow
 - Second Layer Libs:
 - Keras - over TensorFlow
 - Fast.ai - over PyTorch



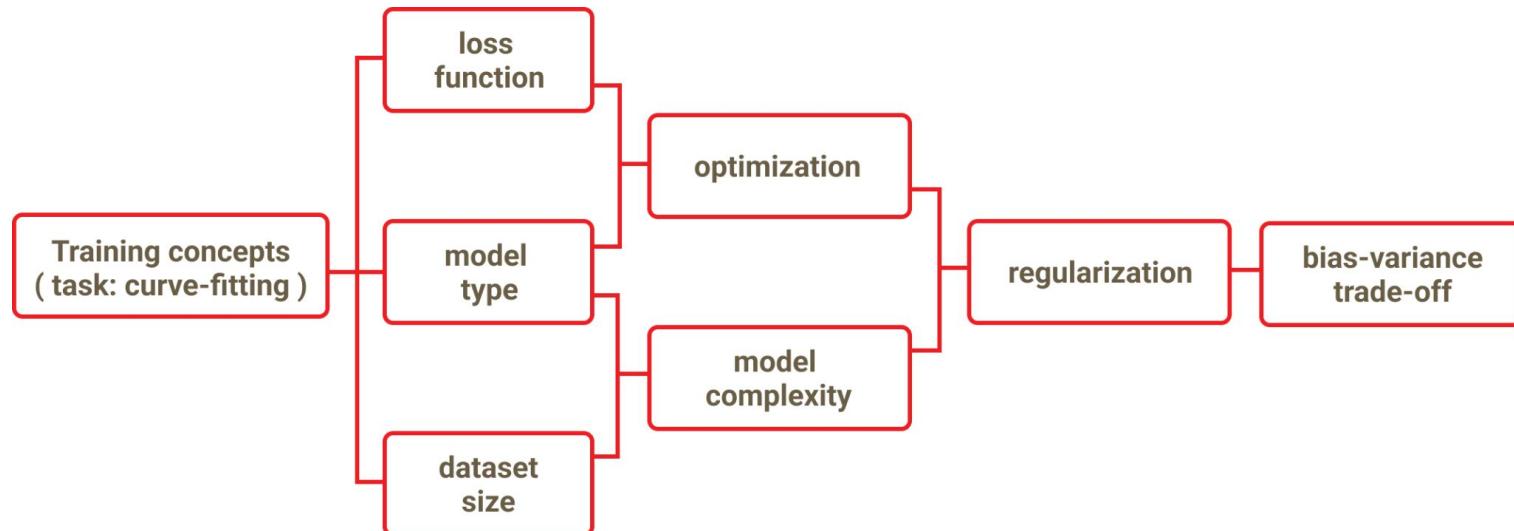
Recap - Deep Learning Motivation



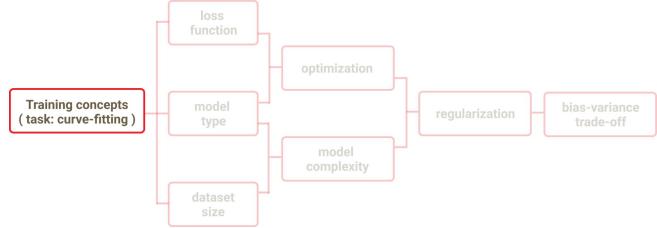
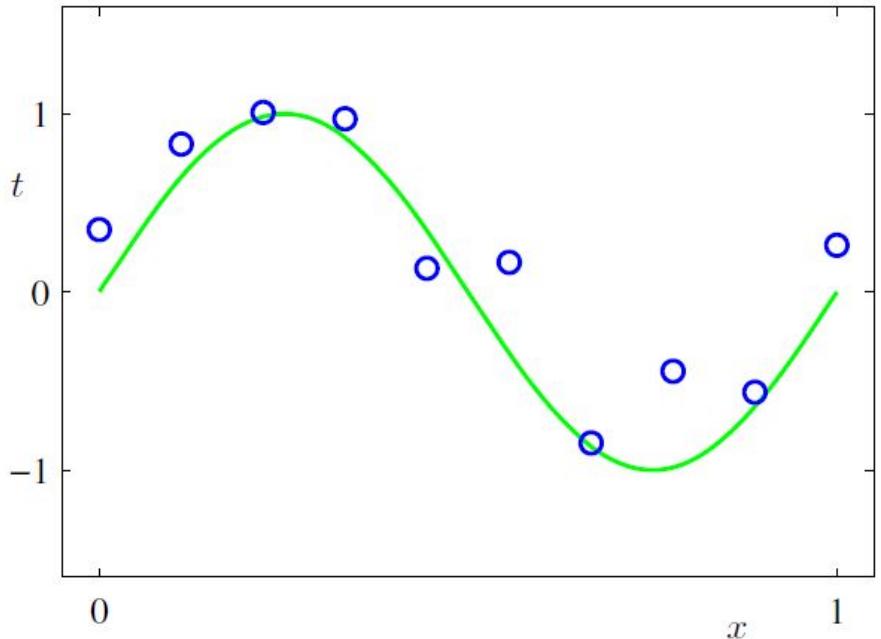


II. Training Concepts

Training Concepts



Task: Curve-Fitting



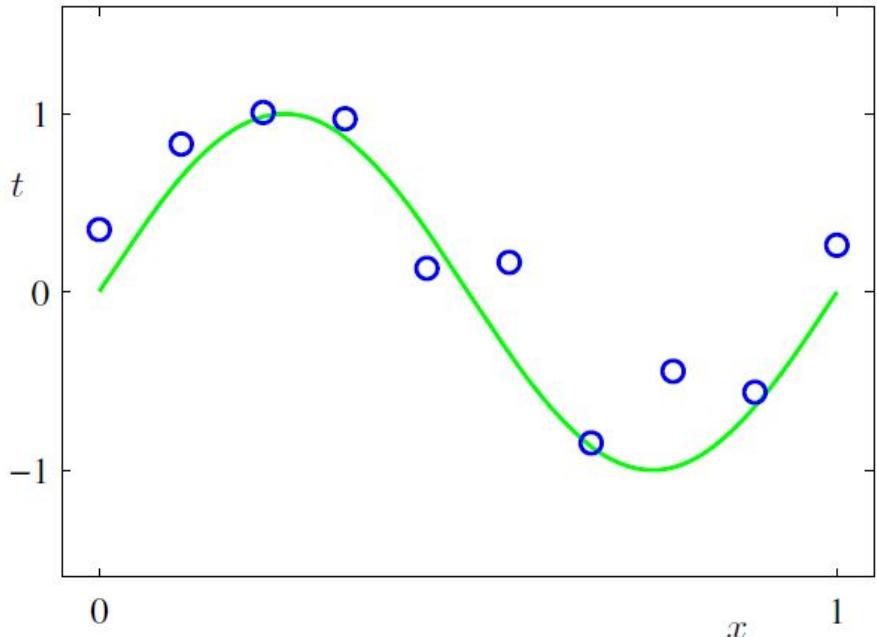
- Training set: $(x_1, t_1), (x_2, t_2), \dots, (x_N, t_N)$
 - $x_i, t_i \in \mathbb{R}$

Data generating process: $\sin(2\pi x) + \text{noise}$

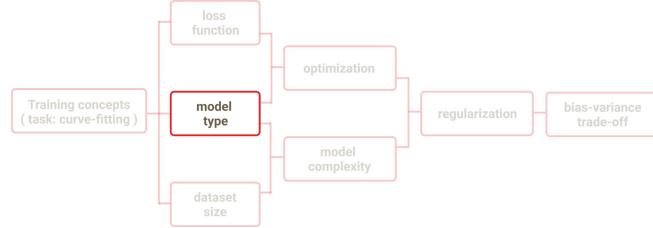
- Underlying regularity which we wish to learn
- Individual observations are corrupted by random noise

- **Goal:** predict \hat{t} for some new \hat{x}
 - regression problem

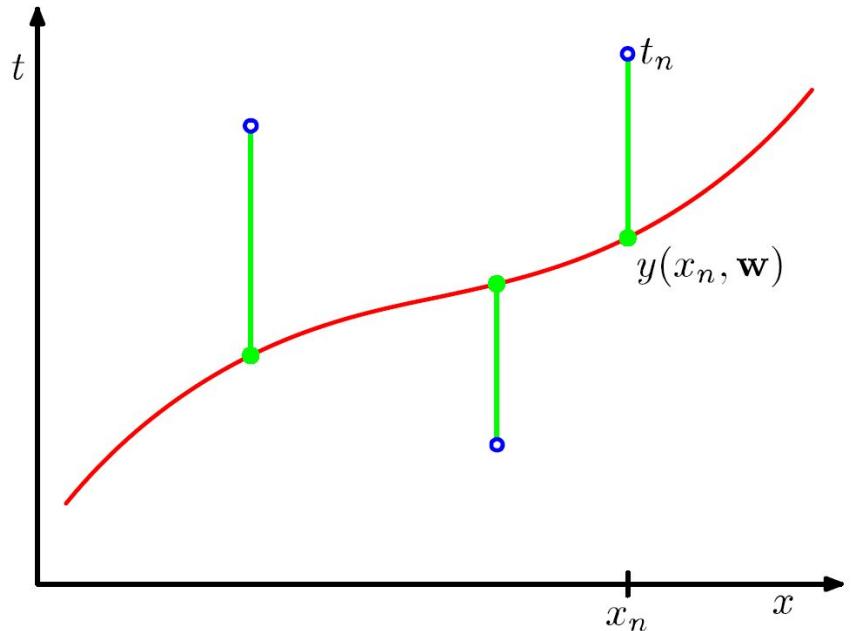
Polynomial Curve Fitting



- Fit the data using a polynomial function
$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M$$
- M - order of the polynomial - model complexity
- \mathbf{w} - parameters to be learned
- y - nonlinear function of x
- y - linear function of the parameters \mathbf{w}

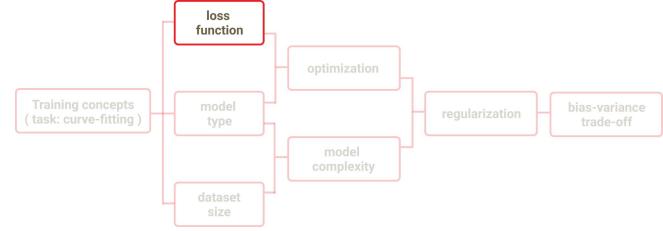


Loss Function



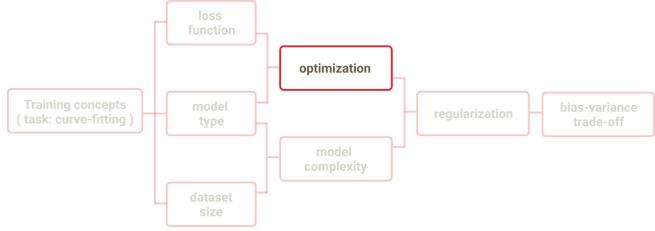
- Sum of squares error

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

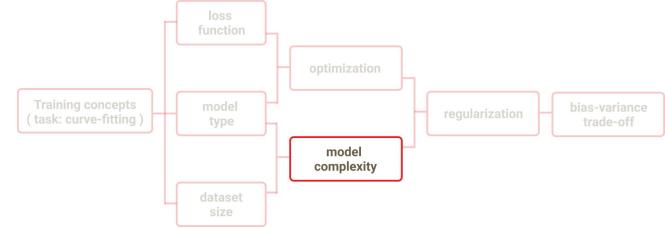
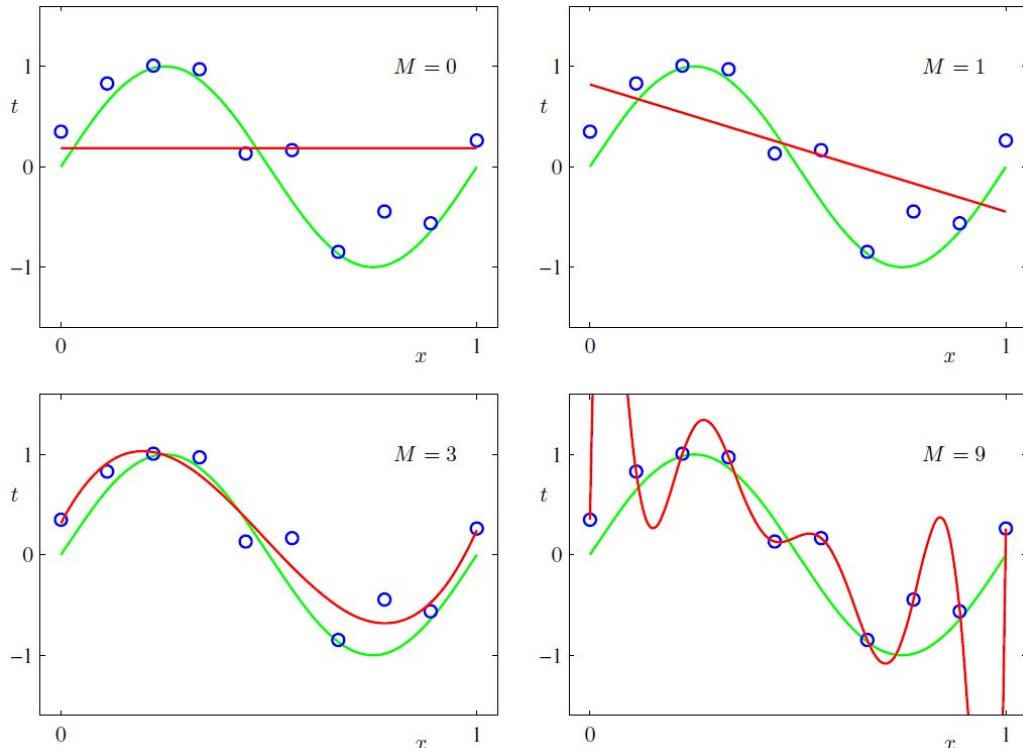


Closed Form Solution

- Goal is to minimize $E(w)$
- $E(w)$ quadratic function of the parameters
 - Derivative w.r.t. parameters is linear in the parameters
 - Unique solution w^* can be found in closed form



Model Complexity



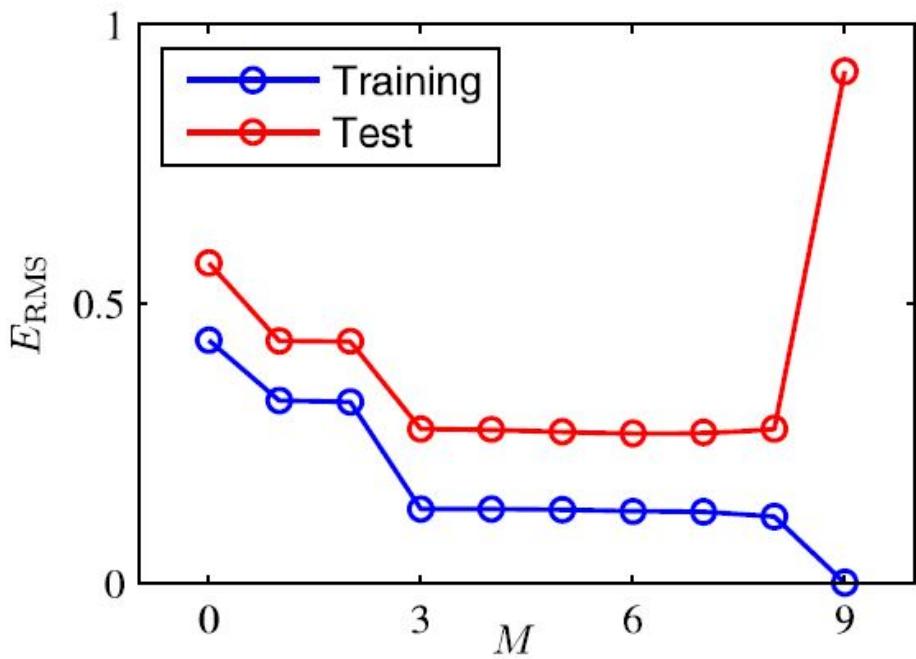
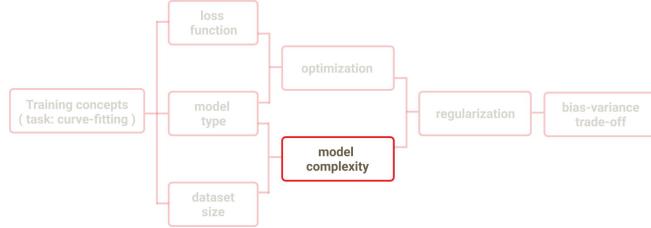
$$y(x, \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M$$

Model Complexity

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

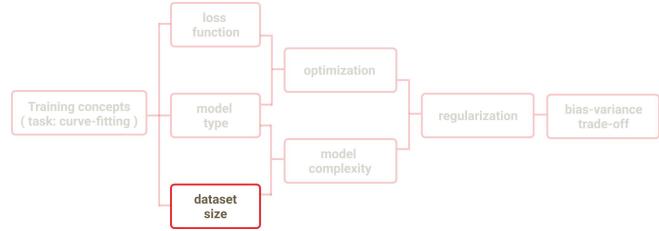
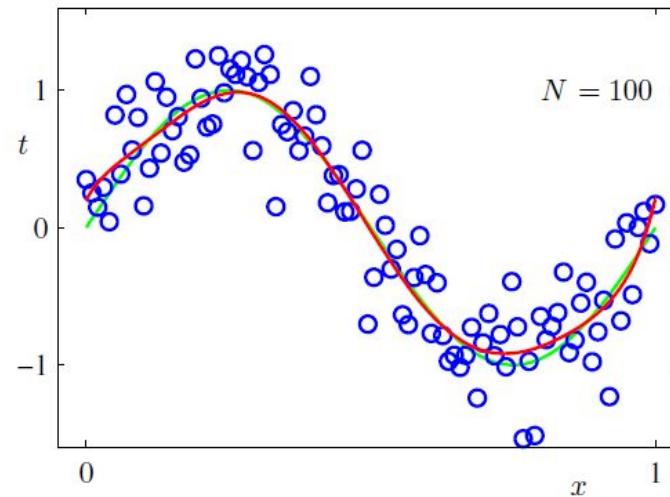
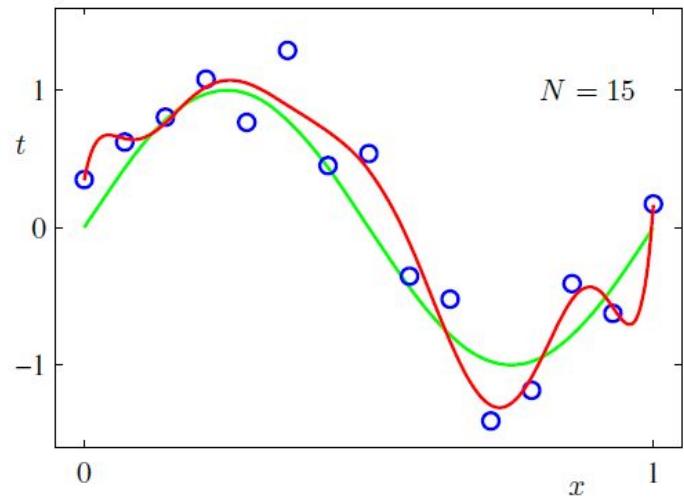
Root-Mean-Square Error:

$$E_{\text{RMS}} = \sqrt{\frac{2E(\mathbf{w}^*)}{N}}$$



Increase training set

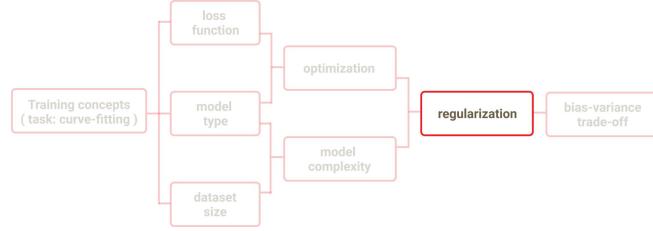
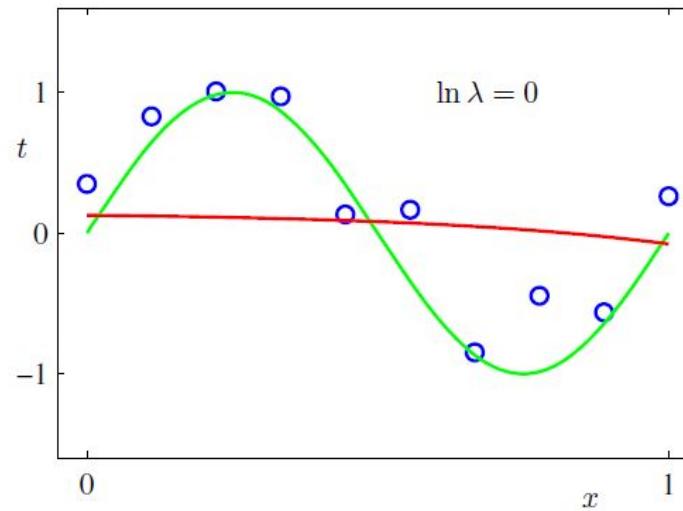
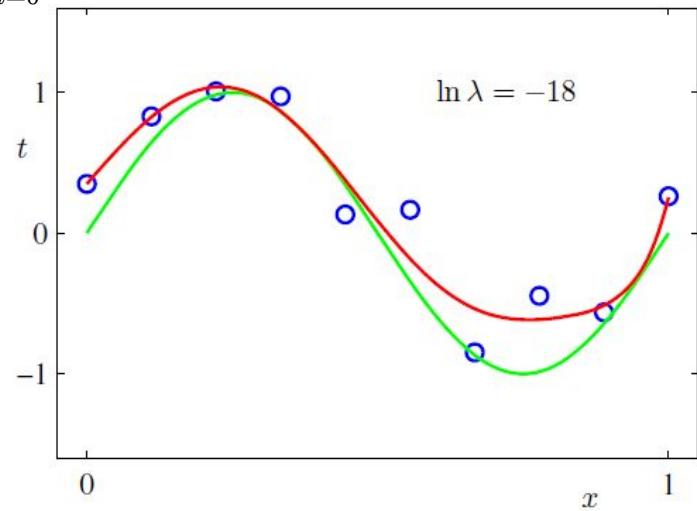
- $M = 9$



Regularization

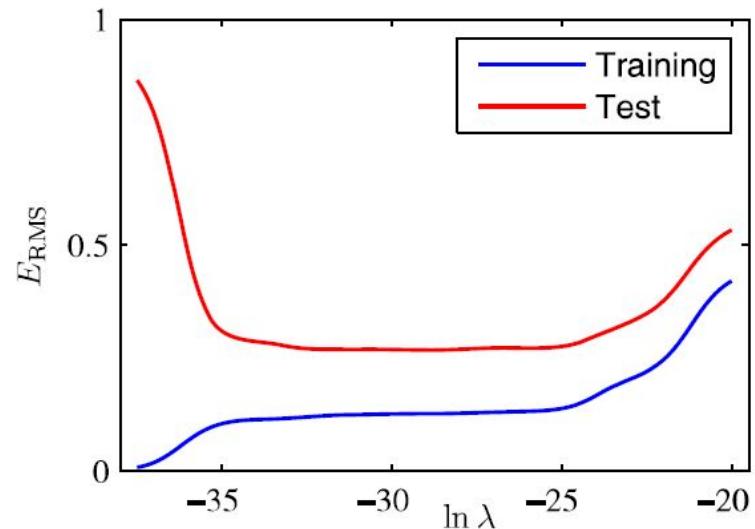
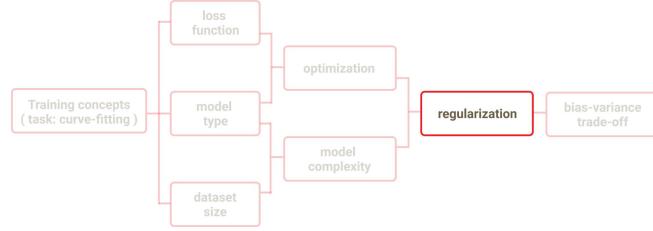
$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$\|\mathbf{w}\|^2 = \sum_{i=0}^M w_i^2$$



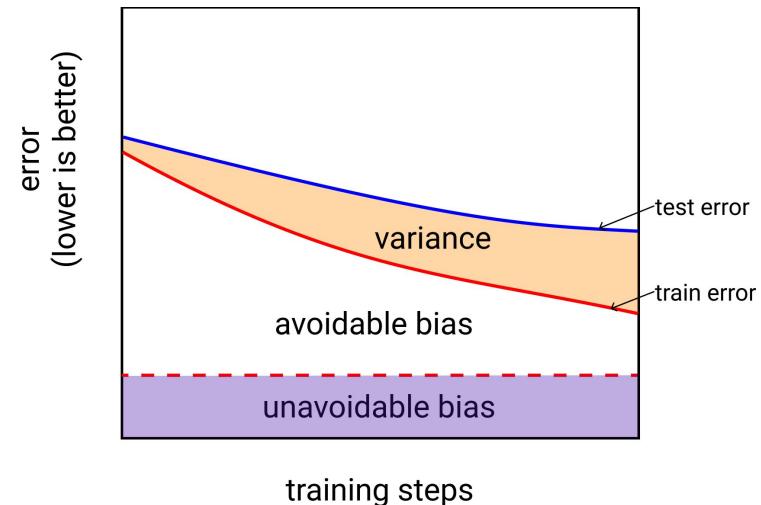
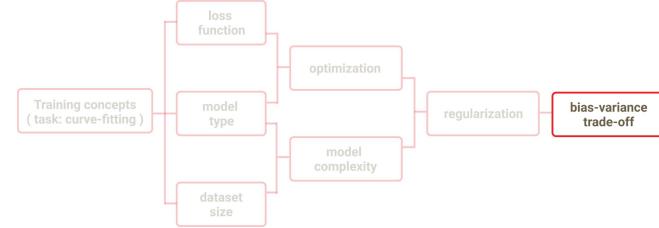
Regularization

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01

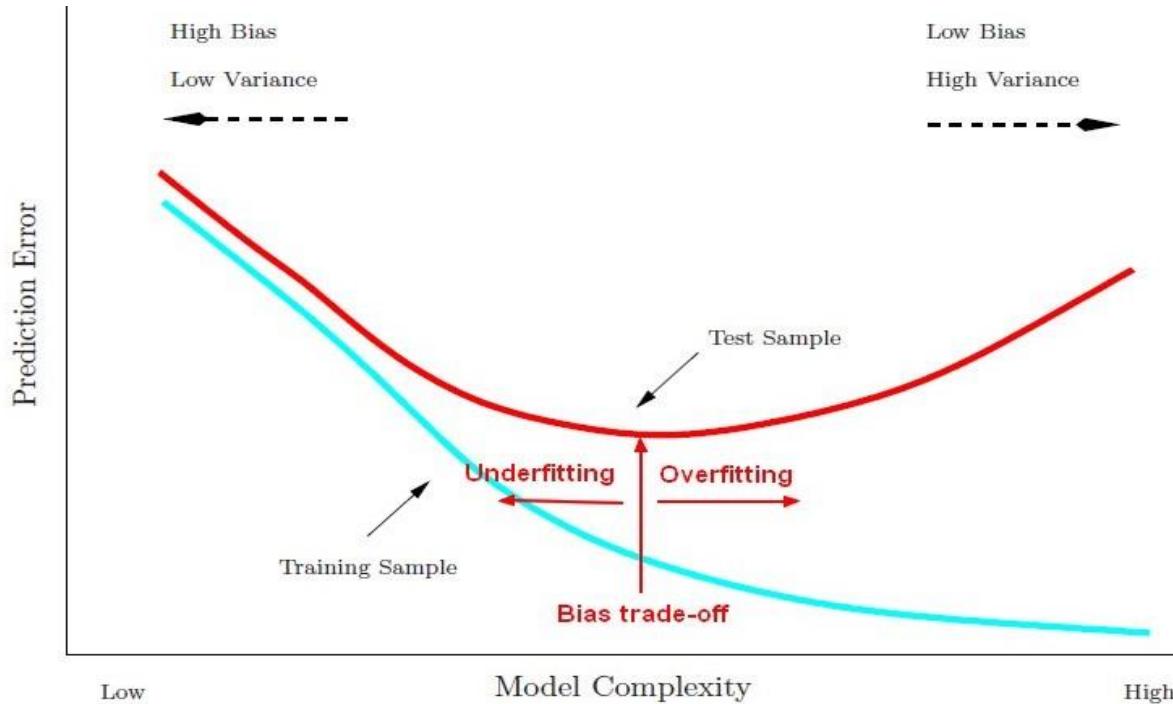
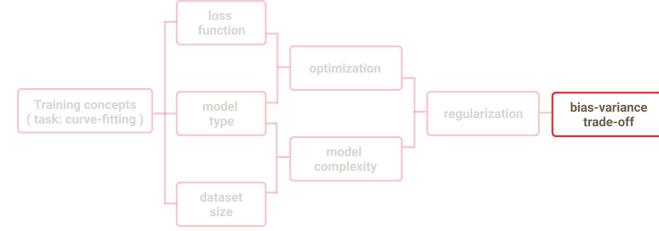


Bias and Variance

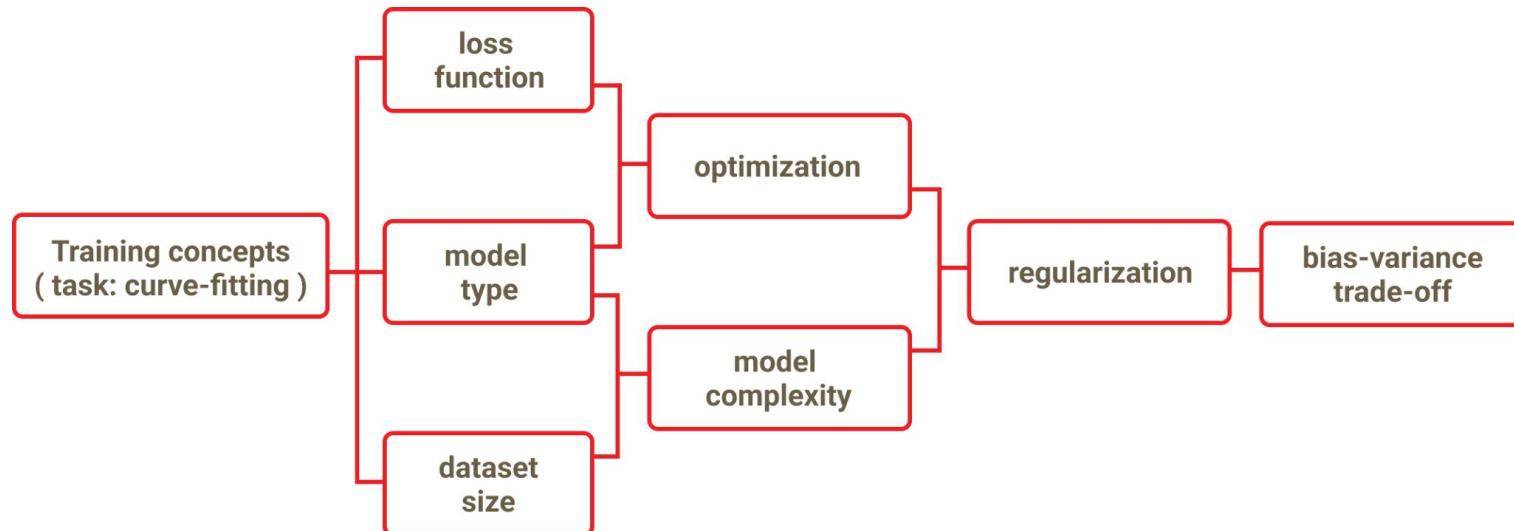
- Algorithm's Bias
 - Error on training set
- Algorithm's Variance
 - Performance decrease between train set and test set
- High Variance => Overfitting
- High Bias => Underfitting
- Bias = Unavoidable Bias + Avoidable Bias



Bias-Variance Trade-Off



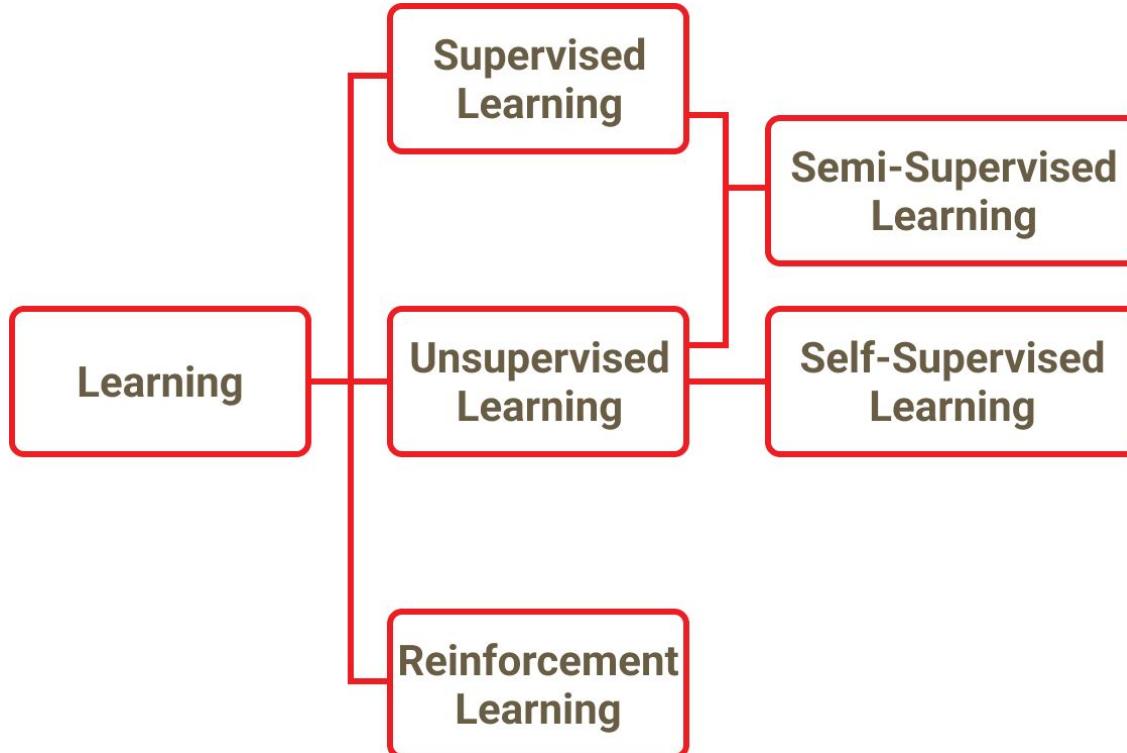
Recap - Training Concepts



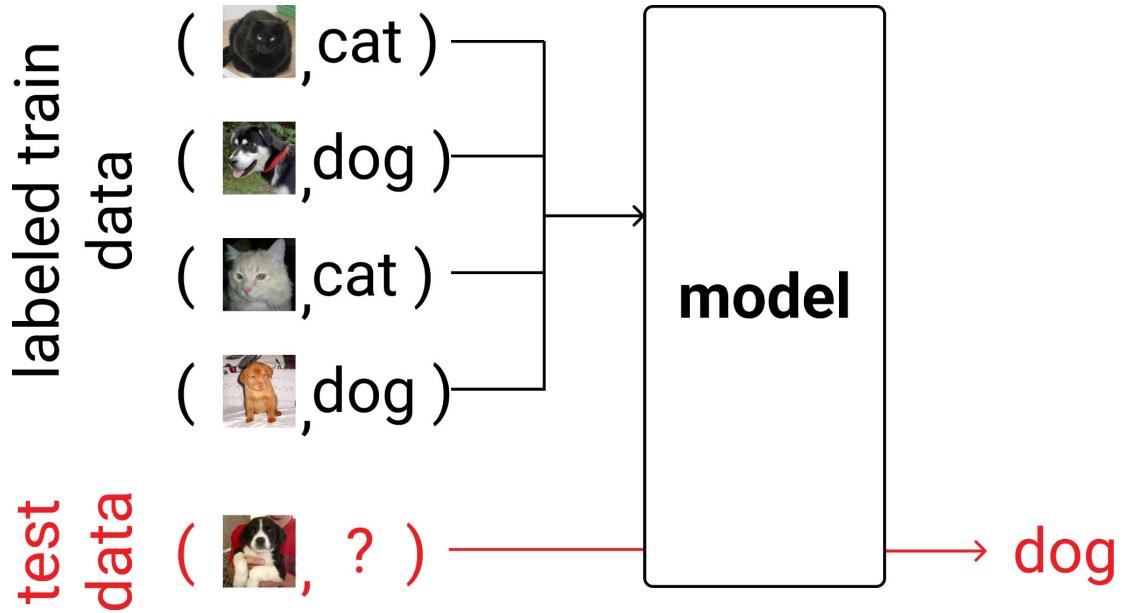
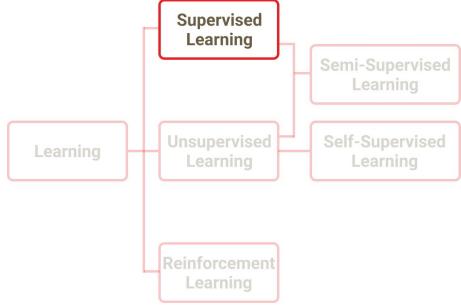


III. Types of Learning

Types of Learning

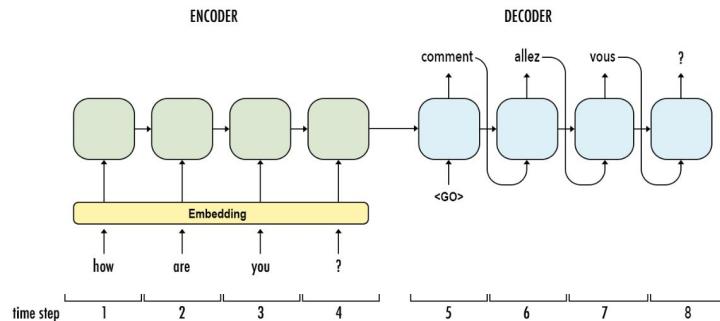
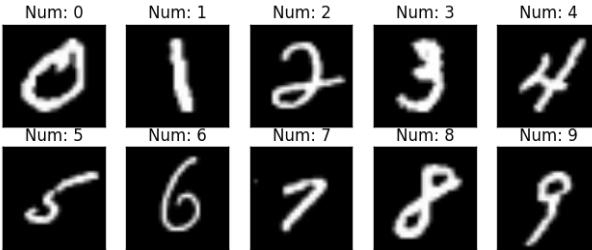
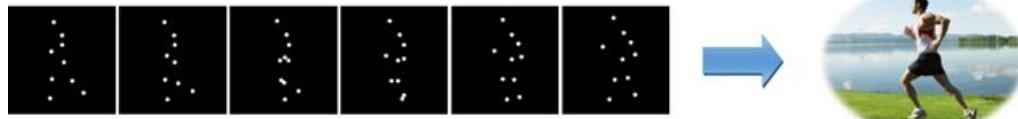
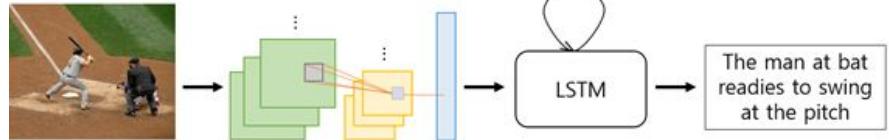


Supervised Learning

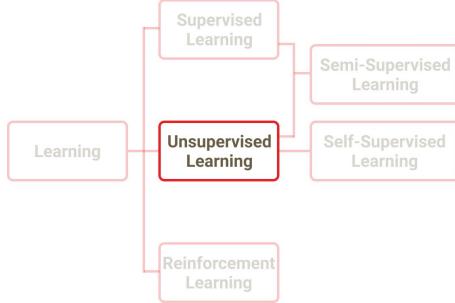
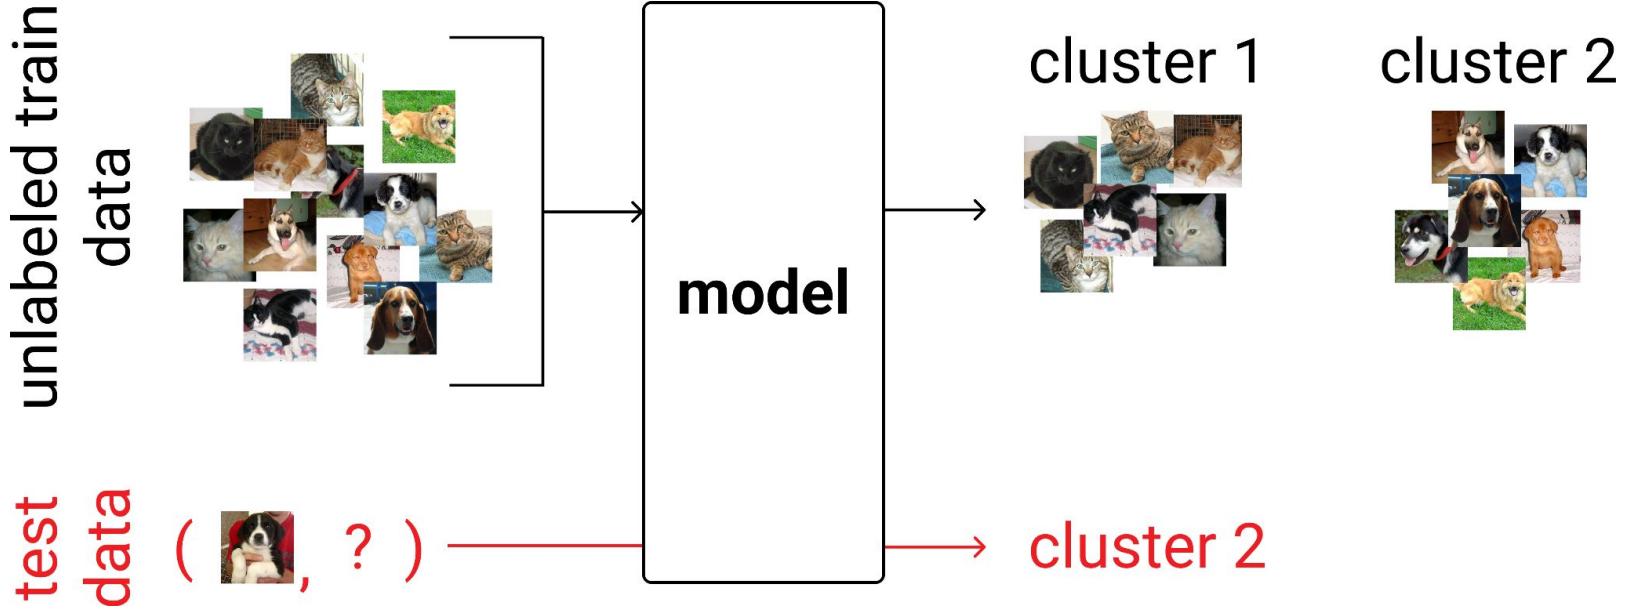


Supervised Learning

- Labeled examples
- Classification or regression tasks

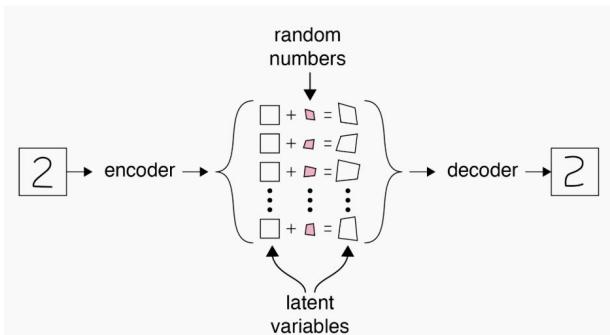
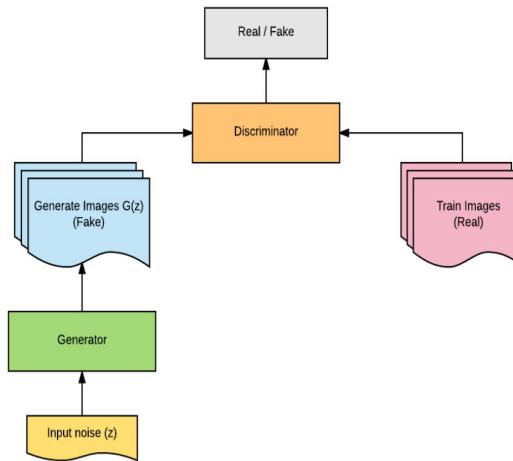
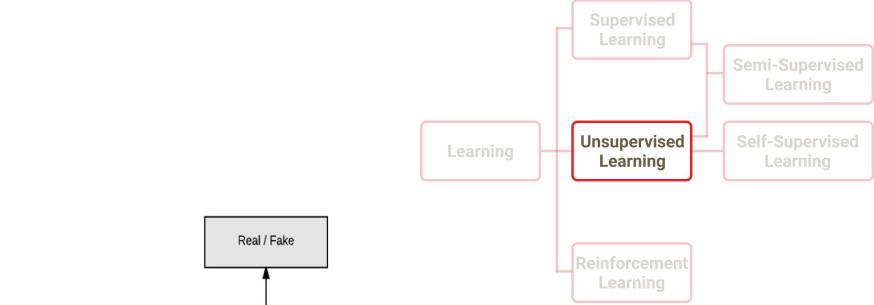
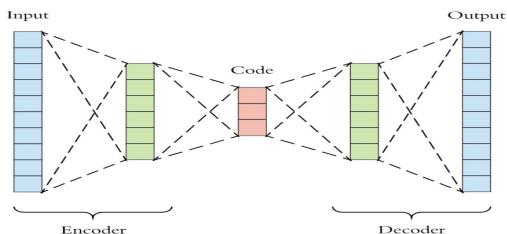
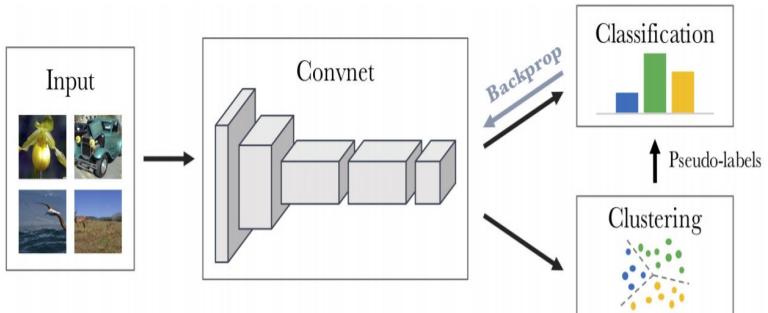


Unsupervised Learning

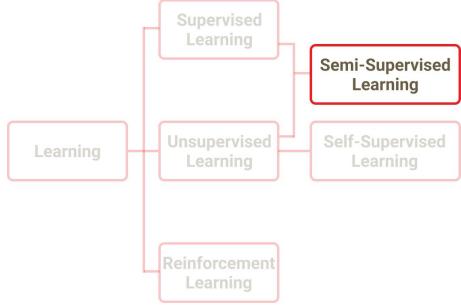
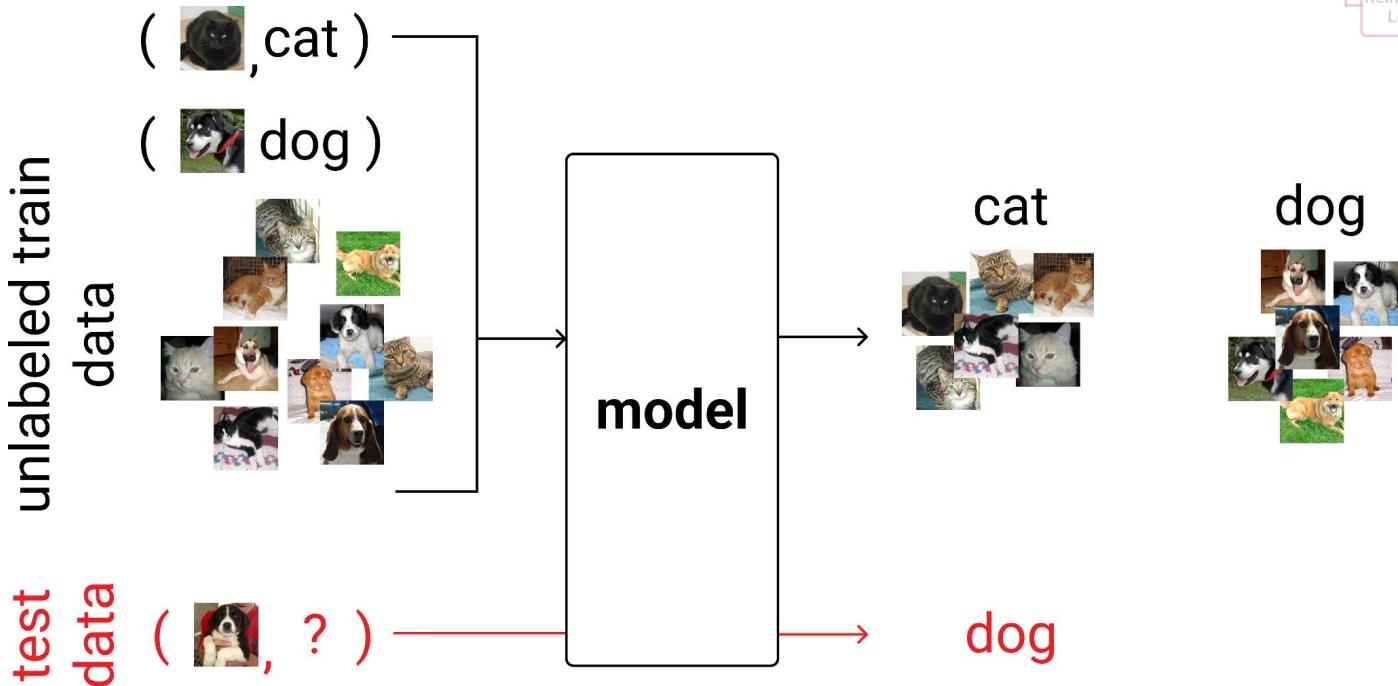


Unsupervised Learning

- Unlabeled data
- Unsupervised features learning

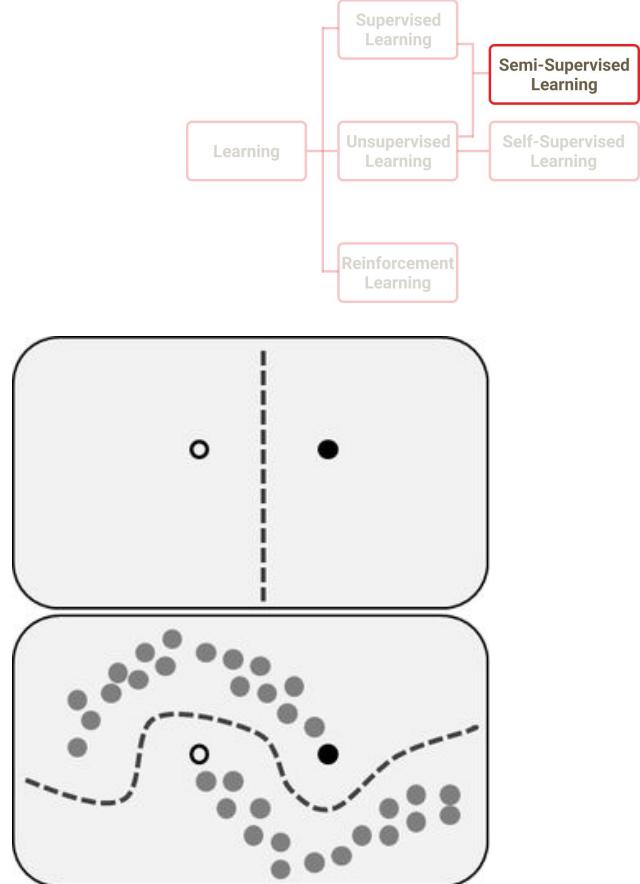


Semi-supervised learning

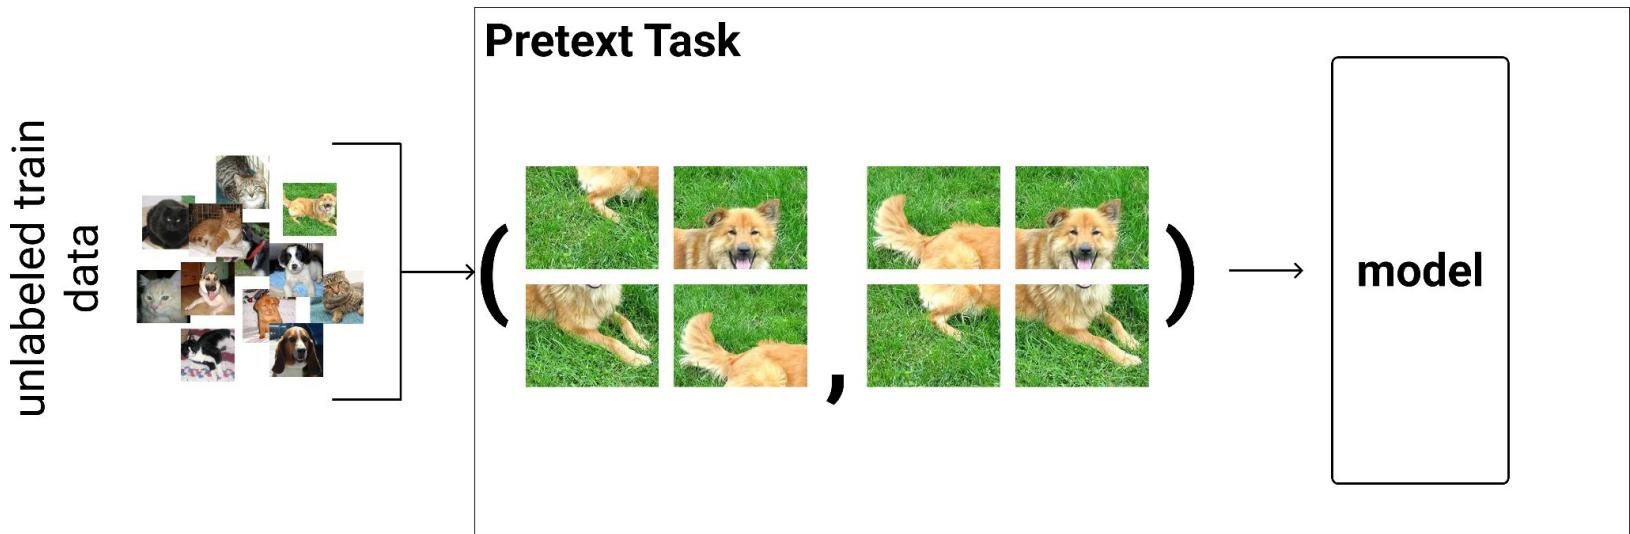
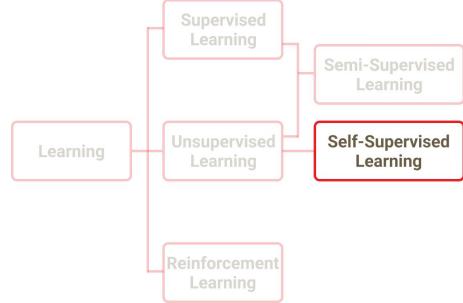


Semi-supervised learning

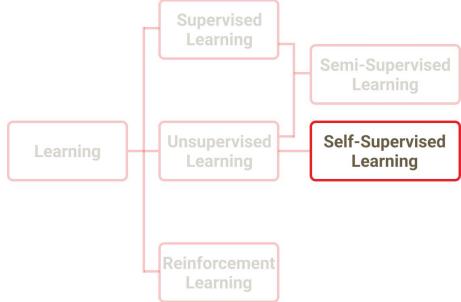
- Small fraction of training examples are labeled
- Tasks:
 - Speech Analysis
 - Webpage Classification
 - Genetic Sequencing



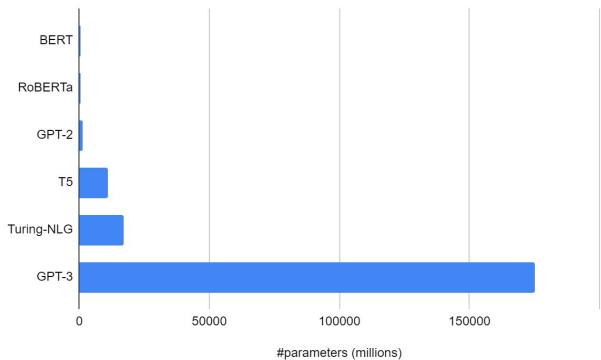
Self-Supervised Learning



Self-Supervised Learning



- GPT-3
 - self-supervised language model - next word prediction
 - 175 billion parameters
 - trained on 10 000 years of continuous speech (at 100 words per minute)



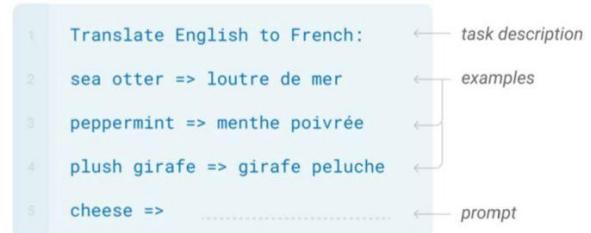
Zero-shot

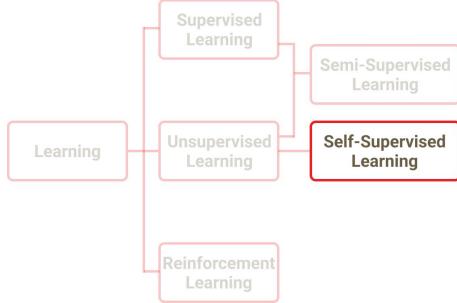
The model predicts the answer given only a natural language description of the task. No gradient updates are performed.



Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.





Self-Supervised Learning

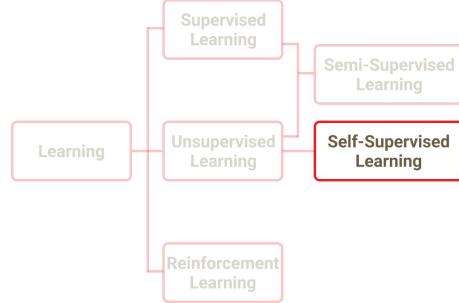
- DALL-E

TEXT PROMPT an armchair in the shape of an avocado....

AI-GENERATED
IMAGES



Self-Supervised Learning



- DALL-E

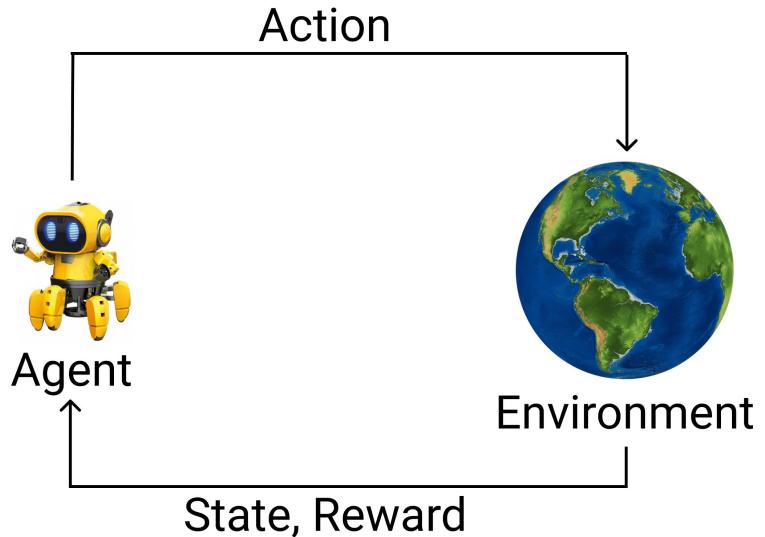
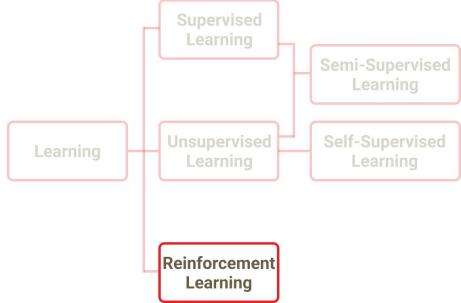
TEXT & IMAGE
PROMPT

the exact same cat on the top as a sketch on the bottom

AI-GENERATED
IMAGES

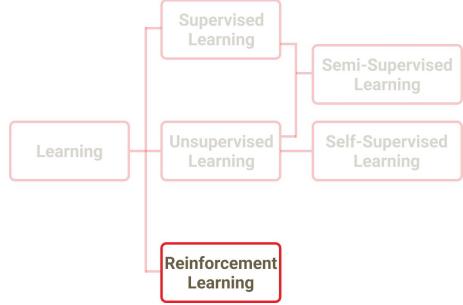


Reinforcement Learning



Reinforcement Learning

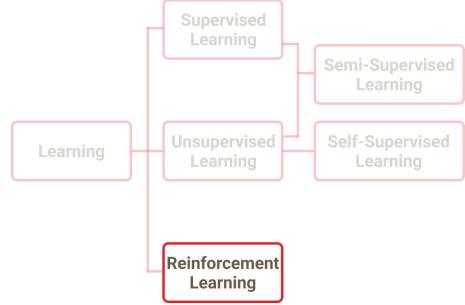
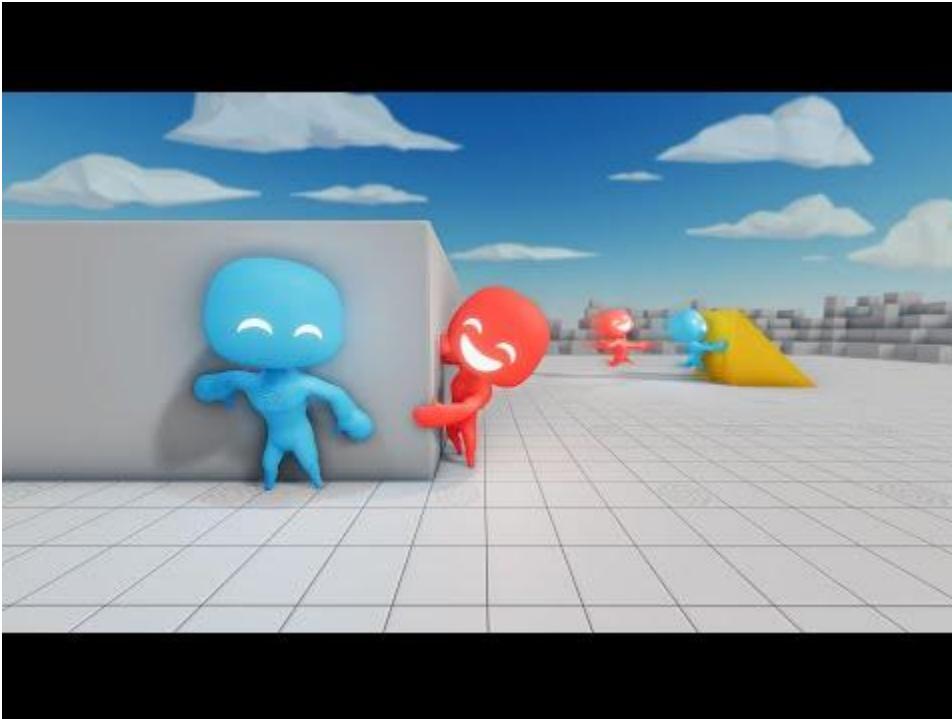
- Optimal behavior for maximal reward
- Applications
 - Robotics
 - Complex games: e.g. AlphaGO
 - Autonomous Driving



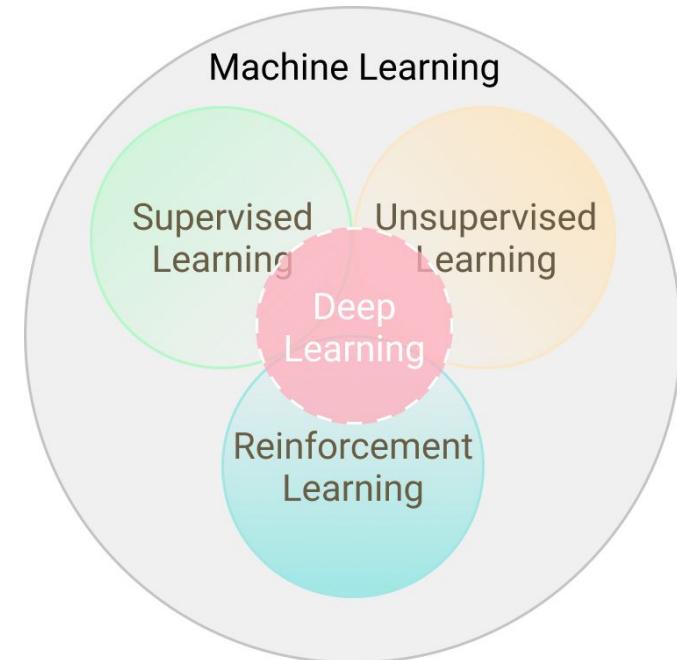
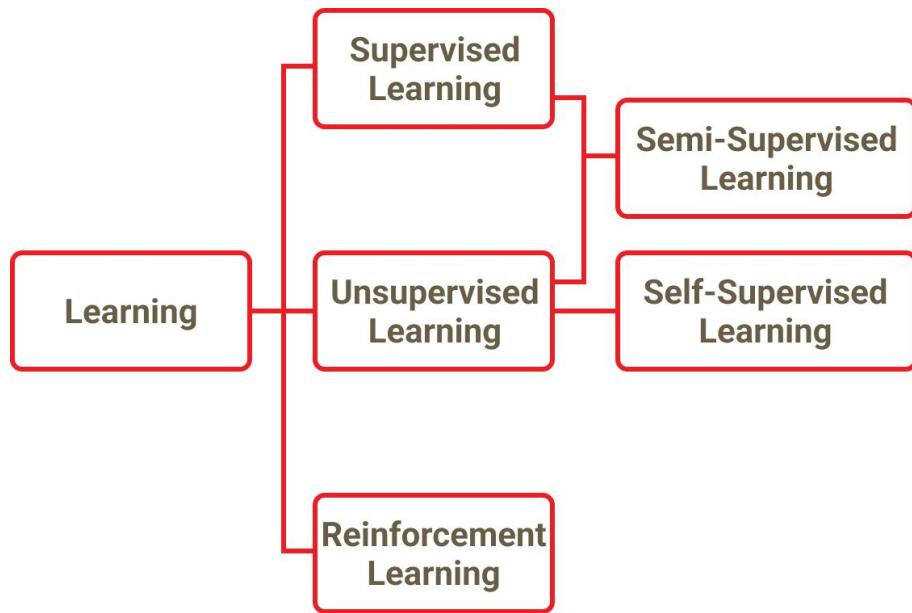
Reinforcement Learning

Multi-Agent

Hide and Seek



Recap - Types of Learning



Next: Neural Networks



Thank you!

References

- Deep Learning - Ian Goodfellow, Yoshua Bengio, Aaron Courville
 - Chapter 1
- Pattern Recognition and Machine Learning - Christopher Bishop
 - Chapter 1
- CS231n: Convolutional Neural Networks for Visual Recognition
 - Module 1